



# LLVM


MI-GEN, [majpetr@fit.cvut.cz](mailto:majpetr@fit.cvut.cz)



# Low Level Virtual Machine

- Started around 2000 at University of Illinois at Urbana-Champaign by Chris Lattner, first public release around 2003
- Now maintained by Apple
  - Default compiler for OS X and IOS
- Users & contributors include: Sony, Adobe, Intel, NVIDIA, XMOS, and many others
- Used in various scenarios – including both AOT and JIT compilers, code analysis, etc.

# Low Level Virtual Machine



and 2000 at University of Illinois at Urbana-Champaign  
by Chris Lattner, first public release

maintained by Apple

compiler for OS X and iOS

Contributors include: Sony, Adobe, Intel, Microsoft, and many others

- Used in various scenarios – including both AOT and JIT compilers, code analysis, etc.



# Low Level Virtual Machine

- Started around 2000 at University of Illinois at Urbana-Champaign by Chris Lattner, first public release around 2003
- Now maintained by Apple
  - Default compiler for OS X and IOS
- Users & contributors include: Sony, Adobe, Intel, NVIDIA, XMOS, and many others
- Used in various scenarios – including both AOT and JIT compilers, code analysis, etc.



# Low Level **V**irtual **M**achine

- Permissive license (BSD-style)
- Modern compiler, written in C++
- Under active development
- Modular design
- Well documented
- Language agnostic
  - Does not care about the frontend
  - Many targets available (x86, ARM, PowerPC, etc.)

# Low Level Virtual Machine

- Permissive license (BSD-style)
- Modern compiler, written in C++
  - Templates & advanced C++ constructs
  - Indecipherable error messages
  - On a plus side, does not use BOOST
- Under active development
- Modular design
- Well documented
- Language agnostic
  - Does not care about the frontend
  - Many targets available (x86, ARM, PowerPC, etc.)

# Low Level Virtual Machine

- Permissive license (BSD-style)
- Modern compiler, written in C++
  - Templates & advanced C++ constructs
  - Indecipherable error messages
  - On a plus side, does not use BOOST
- Under active development
  - Backwards compatibility is frowned upon
- Modular design
- Well documented
- Language agnostic
  - Does not care about the frontend
  - Many targets available (x86, ARM, PowerPC, etc.)



# Low Level Virtual Machine

- Permissive license (BSD-style)
- Modern compiler, written in C++
  - Templates & advanced C++ constructs
  - Indecipherable error messages
  - On a plus side, does not use BOOST
- Under active development
  - Backwards compatibility is frowned upon
- Modular design - kind of
- Well documented
- Language agnostic
  - Does not care about the frontend
  - Many targets available (x86, ARM, PowerPC, etc.)



# Low Level Virtual Machine

- Permissive license (BSD-style)
- Modern compiler, written in C++
  - Templates & advanced C++ constructs
  - Indecipherable error messages
  - On a plus side, does not use BOOST
- Under active development
  - Backwards compatibility is frowned upon
- Modular design - kind of
- ~~Well~~ documented - just better than others, but the bar is low
- Language agnostic
  - Does not care about the frontend
  - Many targets available (x86, ARM, PowerPC, etc.)



# Low Level **V**irtual **M**achine



C

C++

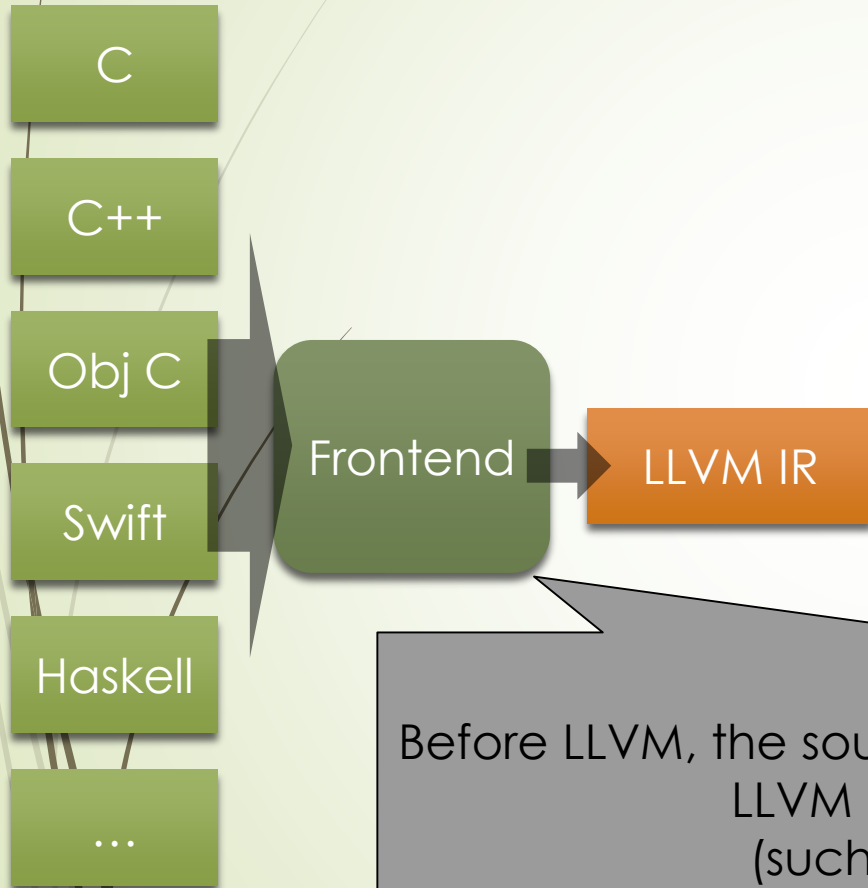
Obj C

Swift

Haskell

...

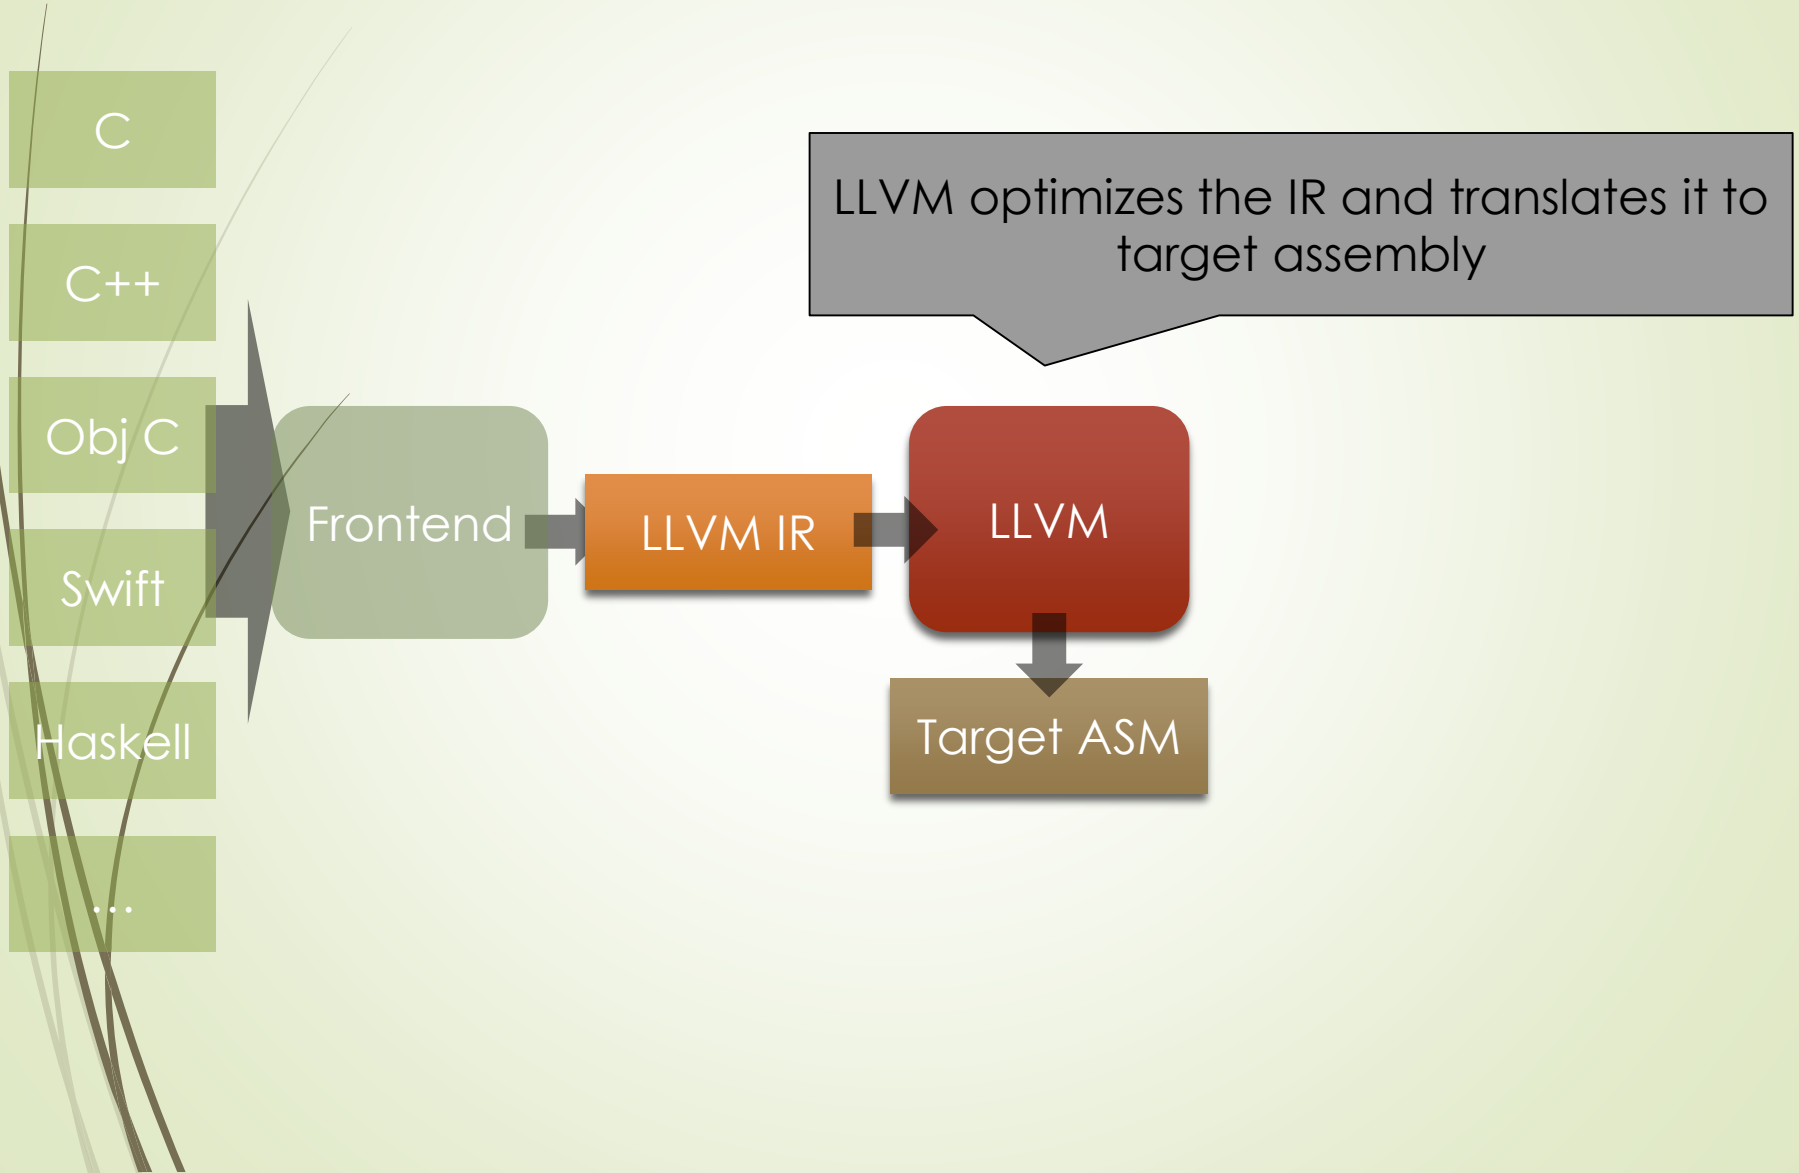
# Low Level Virtual Machine




Before LLVM, the source language must be translated to LLVM IR by custom frontend (such as Clang for C/C++)

In the course of this class, you will create one such frontend

# Low Level Virtual Machine

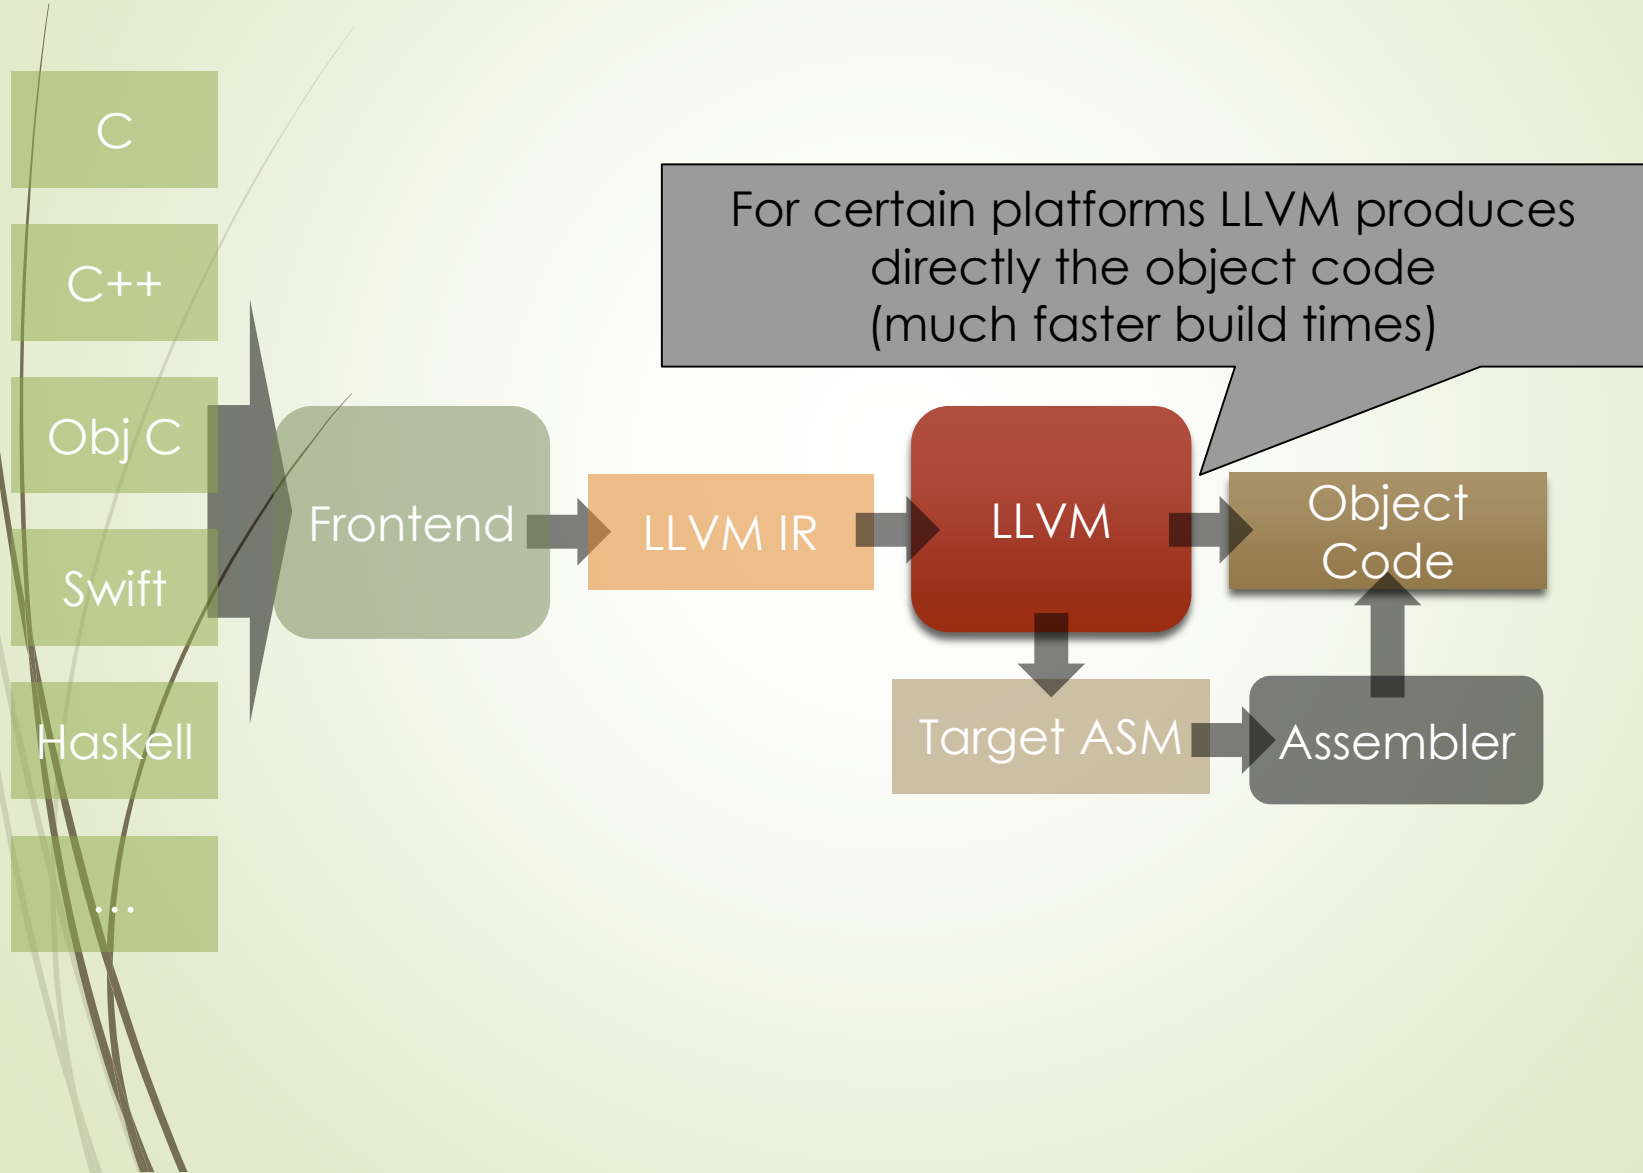


# Low Level Virtual Machine

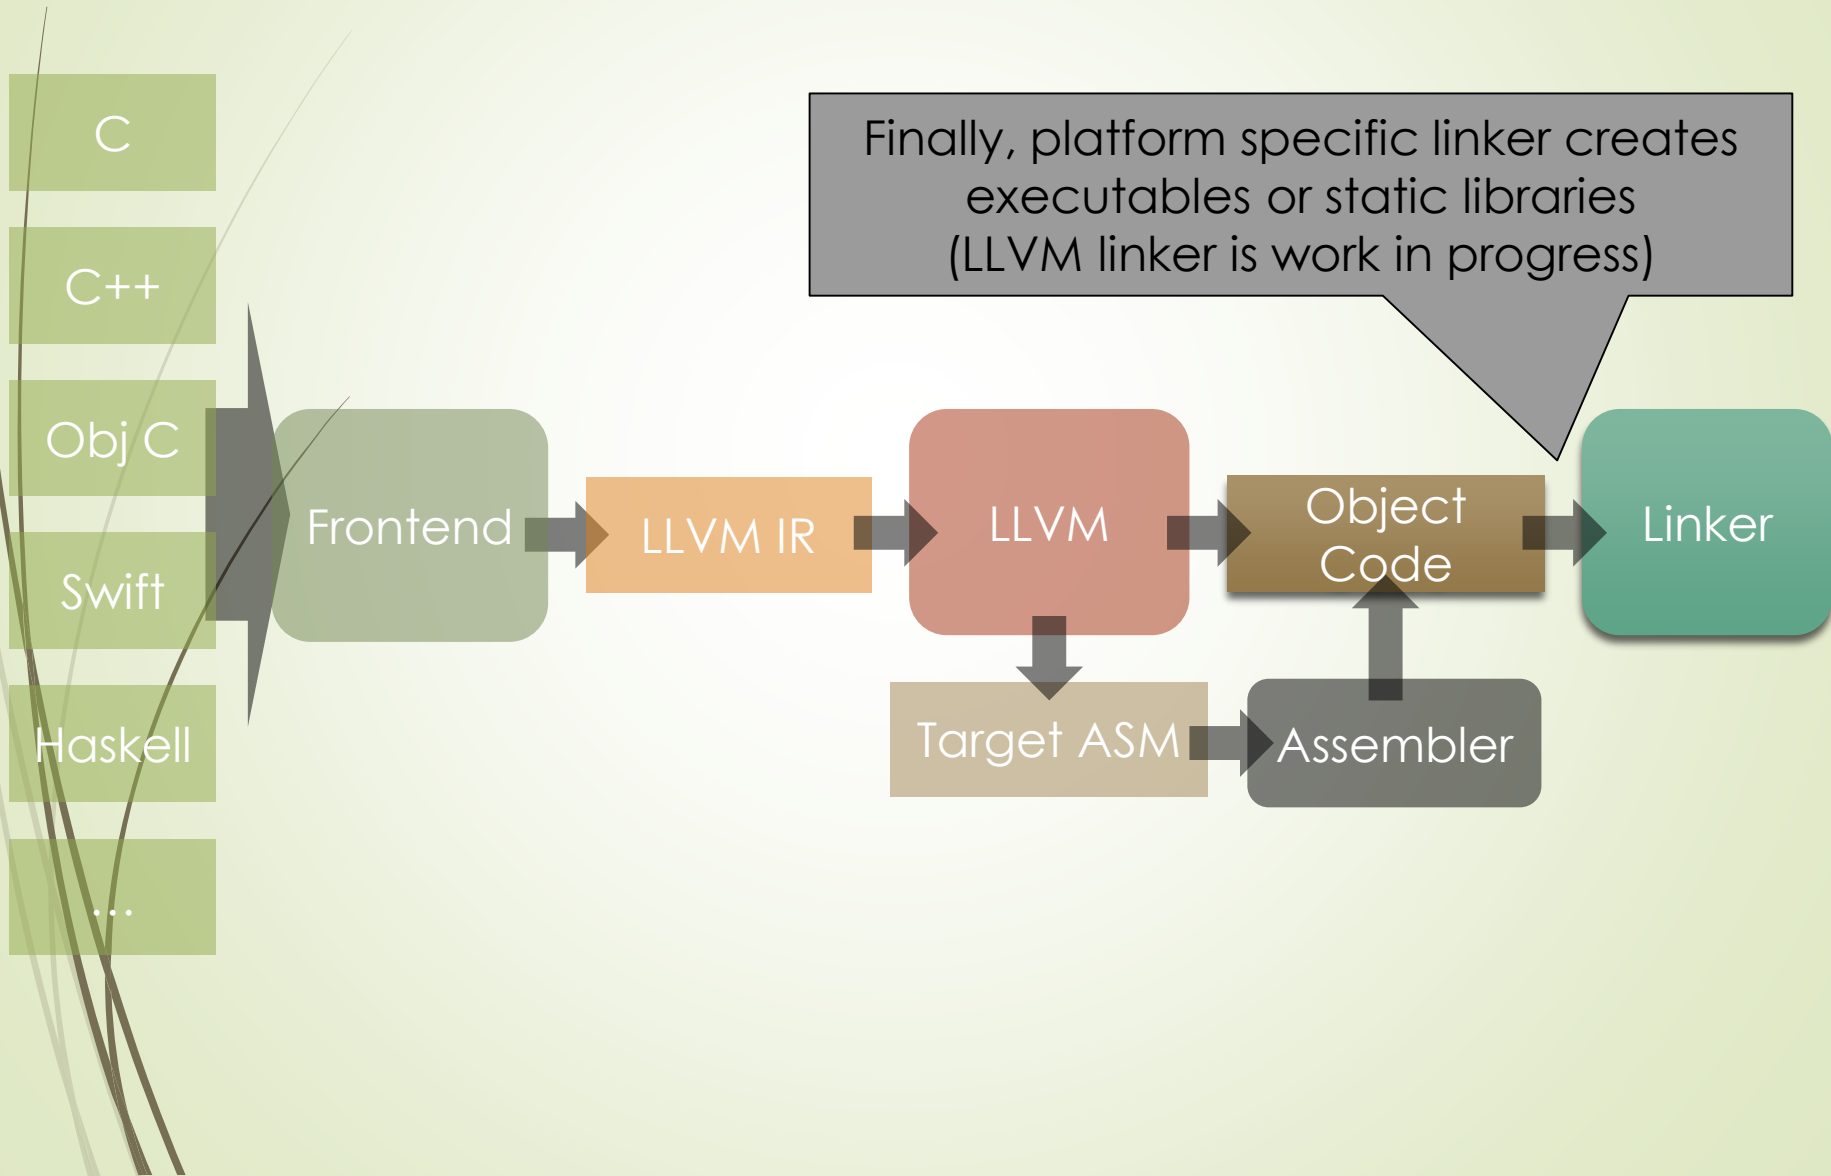


Which is then converted by current platform assembler to object code

# Low Level Virtual Machine

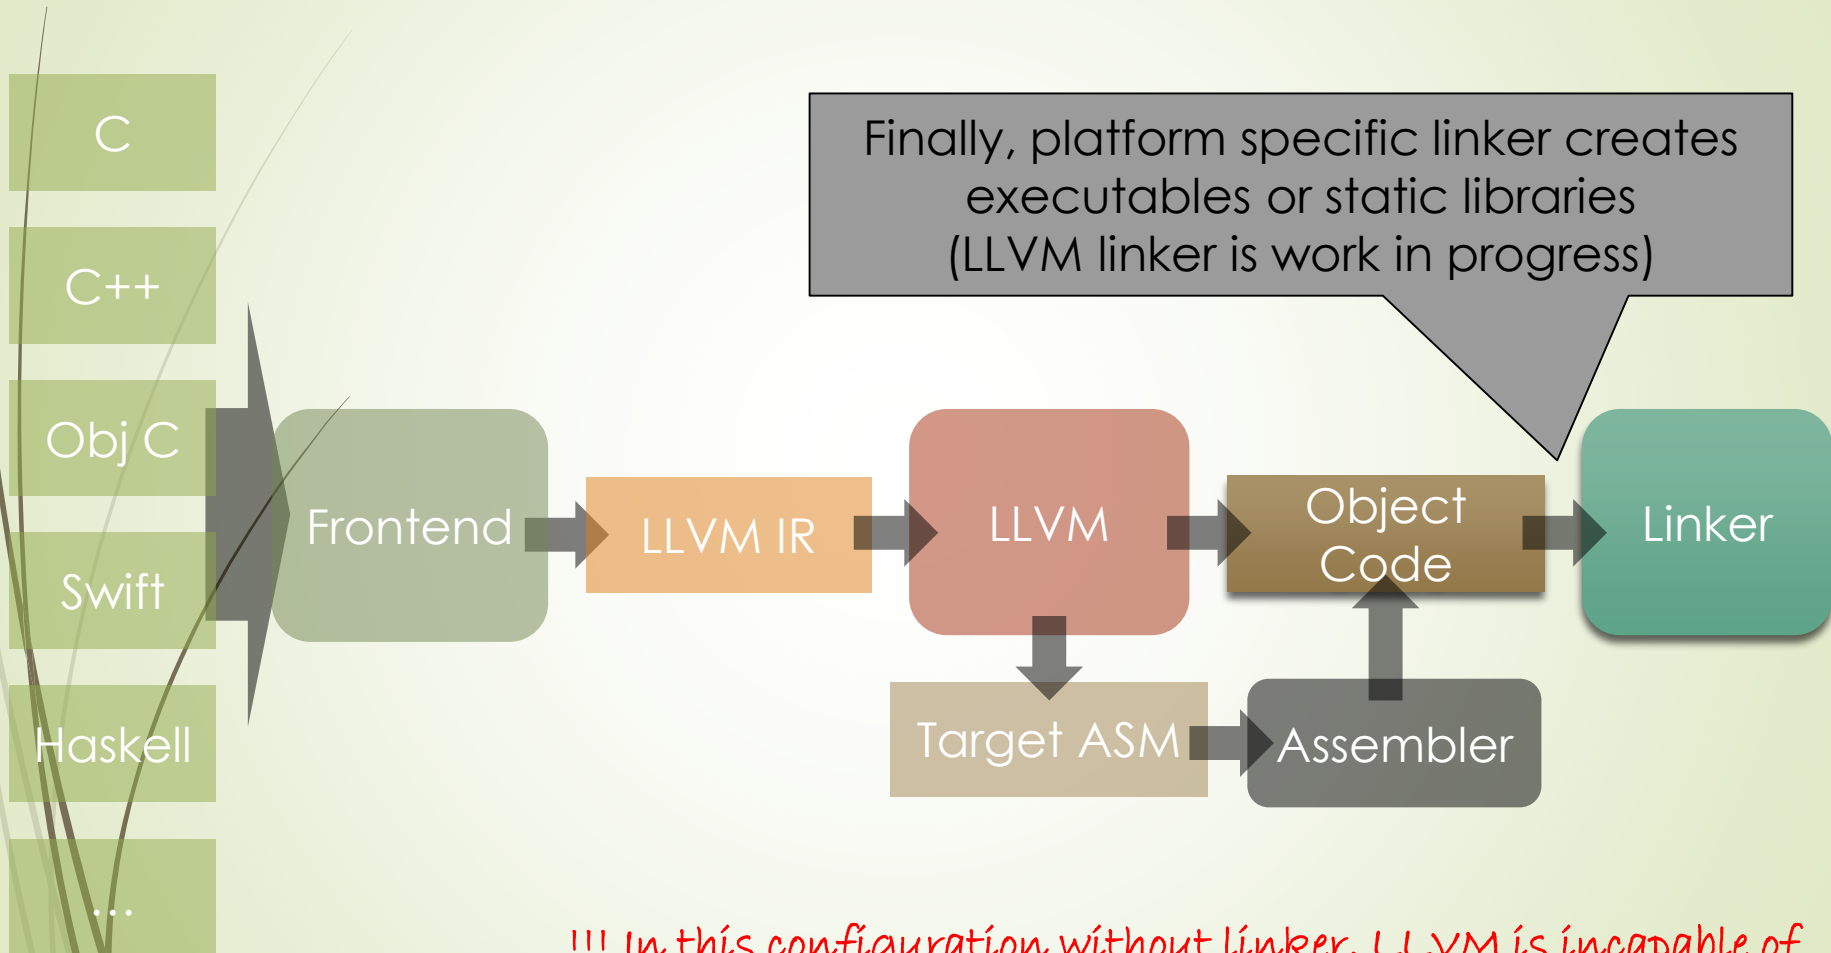


# Low Level Virtual Machine





# Low Level Virtual Machine



!!! In this configuration without linker, LLVM is incapable of link time optimizations !!!

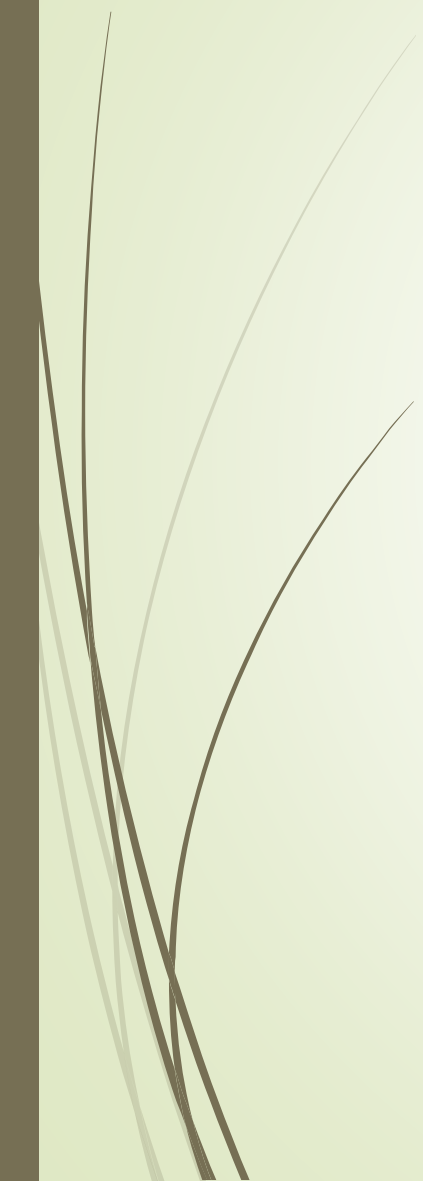


# Low Level **V**irtual **M**achine

- Does not contain frontends
- 



# Low Level **V**irtual **M**achine

- Does not contain frontends
  - Or linker (to be changed soon)
- 



# Low Level **V**irtual **M**achine


- Does not contain frontends
- Or linker (to be changed soon)
- Or assemblers in many platforms
  - Using platform assembler instead



# Low Level Virtual Machine

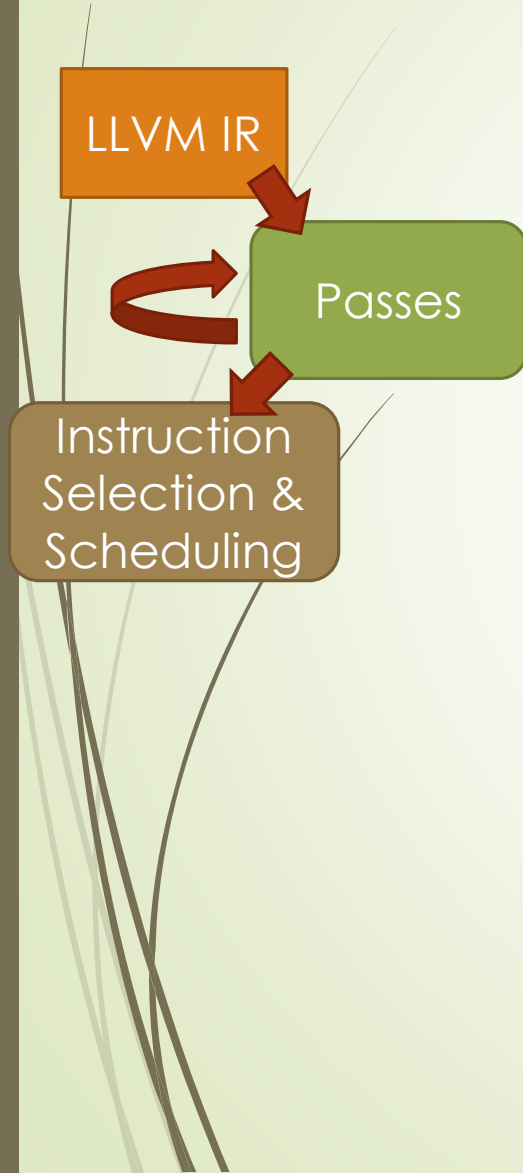
- Does not contain frontends
- Or linker (to be changed soon)
- Or assemblers in many platforms
  - Using platform assembler instead
- LLVM's main focus is on IR optimizations and backend up to assembly
  - SSA IR optimizations
  - Instruction scheduling
  - Register allocation
  - Low level peephole optimizations (to some extent)

# LLVM workflow



- Several analysis (readonly) and optimization passes are performed over the IR
- Passes can run on
  - Modules
  - Functions
  - Basic blocks
- Passes can
  - Depend on other passes
  - Preserve other passes
- LLVM provides a scheduler that calls the passes in a most effective way

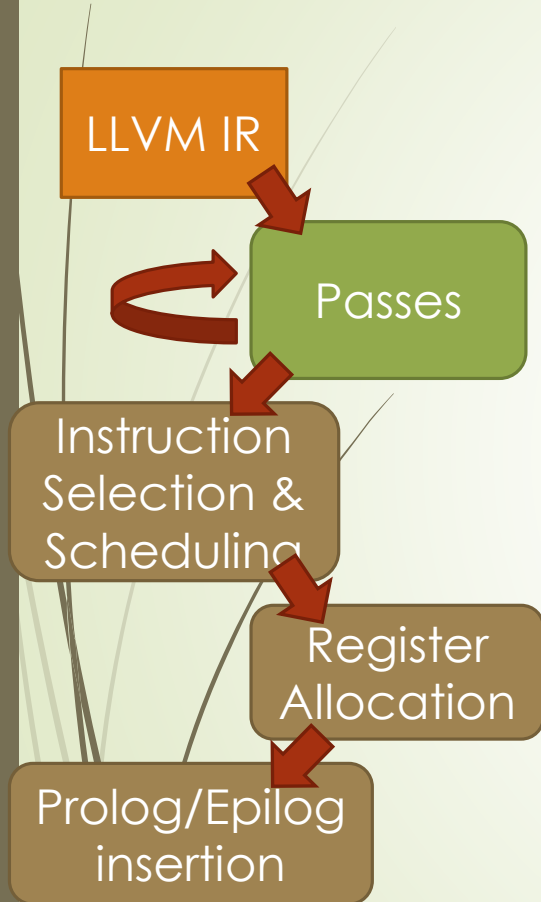
# LLVM workflow



- When IR passes are finished, LLVM enters the target specific backend
- Target specific instructions are selected and attached to the IR
- Afterwards, target specific SSA passes may execute

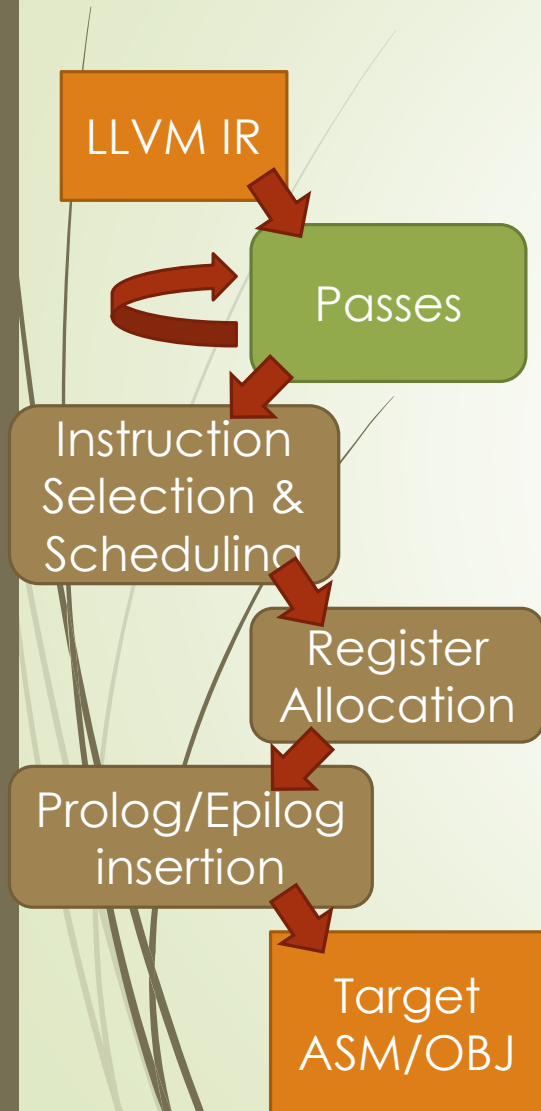


# LLVM workflow



- After SSA target passes, register allocator is executed
  - Linear scan
  - Greedy
  - Region-based
  - GRA
- Function Prologues and epilogues are then inserted
- After these steps, late target optimizations are usually executed
  - peepholer

# LLVM workflow



- Finally IR is dumped as target object, or assembly files



# Bitcode

(LLVM IR)



# Bitcode organization



- Code in LLVM is organized into modules (akin to compilation units)
- Each module contains global variables and functions
  - Unique names
  - Linkage specification
  - Calling convention
  - Visibility,
  - Etc.
- Inside functions, code is organized into basic blocks



# Basic Blocks

- Consecutive sequence of instructions with only one entry and one exit point
- Entry point
  - At the beginning, the only instruction allowed to be jump target
- Exit point
  - At the end, the only instruction allowed to be jump (or return, throw, etc.)
- Once basic block is entered, all its instructions are guaranteed to be executed



# Basic Blocks

```
    <before>  
cond:  
    cmp eax, 0  
    jl next  
    <body>  
    jmp cond  
next:  
    <after>
```

# Basic Blocks


**<before>**  
cond:  
    **cmp** eax, 0  
    **j1** next  
    <body>  
    **jmp** cond  
next:  
    <after>

<before>




# Basic Blocks

```
<before>  
cond:  
    cmp eax, 0  
    jl next  
    <body>  
    jmp cond  
next:  
    <after>
```




# Basic Blocks

```
<before>  
cond:  
    cmp eax, 0  
    jl next  
    <body>  
    jmp cond  
next:  
    <after>
```



# Basic Blocks

```
<before>  
cond:  
    cmp eax, 0  
    jl next  
    <body>  
    jmp cond  
next:  
    <after>
```





# Types

- void (no value, no size)
- integer (variable size in bits)
  - i1, i32, i67837
  - integers do not care about their sign
- Floats
  - half, float, double, ...
- Arrays (of same element types, fixed size)
  - [ 10 x i32], [ 10 x [ 10 x i8]]
- Structures
  - { i32, i32, float, i1 }
  - <{ i8, i8, i32 }> packed structure (padding=0, align=1 byte)
- Functions
  - i32 (i32, i32)
- Pointers (\*)

# Types

- void (no value, no size)
- integer (variable size in bits)
  - i1, i32, i67837
  - integers do not care about their sign
- Floats
  - half, float, double, ...
- Arrays (of same element types, fixed size)
  - [ 10 x i32], [ 10 x [ 10 x i8]]
- Structures
  - { i32, i32, float, i1 }
  - <{ i8, i8, i32 }> packed structure (padding=0, align=1 byte)
- Functions
  - i32 (i32, i32)
- Pointers (\*)

+ fp128  
+ fp80 (x86)  
+ ppc 128bit fp  
+ mmx types  
+ vectors (SIMD)  
+ labels, tokens, metadata, ...



# Variables

- ▶ LLVM provides three types of variables
- ▶ Stack allocated variables
- ▶ Global variables (prefixed with @)
- ▶ Local variables (prefixed with %)



# Variables



- ▶ LLVM provides three types of variables
- ▶ Stack allocated variables
  - ▶ `alloca` = creates variable on stack and returns a pointer to it
  - ▶ `load` & `store`
- ▶ Global variables (prefixed with `@`)
- ▶ Local variables (prefixed with `%`)





# Variables

- ▶ LLVM provides three types of variables
- ▶ Stack allocated variables
- ▶ Global variables (prefixed with @)
  - ▶ Contain pointers to the global variables (i.e. similar to stack allocated variables)
- ▶ Local variables (prefixed with %)



# Variables

- ▶ LLVM provides three types of variables
- ▶ Stack allocated variables
- ▶ Global variables (prefixed with @)
- ▶ Local variables (prefixed with %)
  - ▶ Result of each instruction goes into new local variable
  - ▶ In fact, the variable and the instruction are the same thing in LLVM
  - ▶ This makes the LLVM variables SSA values



# Variables

- ▶ LLVM provides three types of variables
- ▶ Stack allocated variables
- ▶ Global variables (prefixed with @)
- ▶ Local variables (prefixed with %)
  - ▶ Result of each instruction goes into new local variable
  - ▶ In fact, the variable and the instruction are the same thing in LLVM
  - ▶ This makes the LLVM variables SSA values
  - ▶ *Because global variables are only pointers to variable locations, they are SSA values too*



# Single Static Assignment

- Each variable is defined exactly once
- Each variable's definition dominates all its uses

(Cytron, Ferrante, Rosen, Wegman, Zadeck, IBM & Brown 1989)

# Single Static Assignment

- Each variable is defined exactly once
- Each variable's definition dominates all its uses

(Cytron, Ferrante, Rosen, Wegman, Zadeck, IBM & Brown 1989)

- Over its lifetime, a variable is replaced by its versions, new version created whenever the variable is assigned:

$x = 1$

$x = 2$

$z = x$

->

$x1 = 1$

$x2 = 2$

$z1 = x2$

# Single Static Assignment

- Each variable is defined exactly once
- Each variable's definition dominates all its uses

(Cytron, Ferrante, Rosen, Wegman, Zadeck, IBM & Brown 1989)

- Over its lifetime, a variable is replaced by its versions, new version created whenever the variable is assigned:

$x = 1$		$x1 = 1$
$x = 2$	$\rightarrow$	$x2 = 2$
$z = x$		$z1 = x2$

This nicely simplifies the use-def chains (so we can trivially see that  $x1$  is never used and thus eliminate the assignment)



# PHI nodes

- SSA is simple when we have only sequential code, but what if a value may come from different places:

```
if (...)  
    a = 2;  
else  
    a = 3;  
b = a
```

# PHI nodes

- SSA is simple when we have only sequential code, but what if a value may come from different places:

```
if (...)  
    a = 2;  
else  
    a = 3;  
b = a
```

->

```
if (...)  
    a1 = 2;  
else  
    a2 = 3;  
b1 = ???
```



# PHI nodes

- SSA is simple when we have only sequential code, but what if a value may come from different places:

```
if (...)  
    a = 2;  
else  
    a = 3;  
b = a
```

->

```
if (...)  
    a1 = 2;  
else  
    a2 = 3;  
b1 = ???
```

- PHI nodes generate new values by selecting “by magic” the proper values based on where the control flow arrived from

# PHI nodes

- SSA is simple when we have only sequential code, but what if a value may come from different places:

```
if (...)  
    a = 2;  
else  
    a = 3;  
b = a
```

->

```
if (...)  
    a1 = 2;  
else  
    a2 = 3;  
a3 = phi(a1,a2)  
b1 = a3
```

- PHI nodes generate new values by selecting “by magic” the proper values based on where the control flow arrived from



# Bitcode Instructions

- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)



# Bitcode Instructions

- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)



# Bitcode Instructions

- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic



# Bitcode Instructions

- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic
- Vector operations
  - Extract, insert, shuffle



# Bitcode Instructions

- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic
- Vector operations
  - Extract, insert, shuffle
- Memory access & addressing
  - allocation, load, store, fences, getelementptr



# Bitcode Instructions

- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic
- Vector operations
  - Extract, insert, shuffle
- Memory access & addressing
  - allocation, load, store, fences, getelementptr
- Conversions
  - Zext, sext, ...





# Bitcode Instructions

- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic
- Vector operations
  - Extract, insert, shuffle
- Memory access & addressing
  - allocation, load, store, fences, getelementptr
- Conversions
  - Zext, sext, ...
- Calling



# Bitcode Instructions



- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic
- Vector operations
  - Extract, insert, shuffle
- Memory access & addressing
  - allocation, load, store, fences, getelementptr
- Conversions
  - Zext, sext, ...
- Calling
- Comparisons
  - icmp, fcmp, ...



# Bitcode Instructions



- ▀ Terminator instructions
  - ▀ Terminate basic block (jumps, returns, exceptions)
- ▀ Binary Operations
  - ▀ Different opcodes for floats (fadd,...) and integers (add...)
- ▀ Shifts, rotations
  - ▀ arithmetic, logic
- ▀ Vector operations
  - ▀ Extract, insert, shuffle
- ▀ Memory access & addressing
  - ▀ allocation, load, store, fences, getelementptr
- ▀ Conversions
  - ▀ Zext, sext, ...
- ▀ Calling
- ▀ Comparisons
  - ▀ icmp, fcmp, ...
- ▀ Exceptions
  - ▀ Catchpads, landingpads, etc.



# Bitcode Instructions



- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic
- Vector operations
  - Extract, insert, shuffle
- Memory access & addressing
  - allocation, load, store, fences, getelementptr
- Conversions
  - Zext, sext, ...
- Calling
- Comparisons
  - icmp, fcmp, ...
- Exceptions
  - Catchpads, landingpads, etc.
- Virtuals (PHI node)



# Bitcode Instructions



- Terminator instructions
  - Terminate basic block (jumps, returns, exceptions)
- Binary Operations
  - Different opcodes for floats (fadd,...) and integers (add...)
- Shifts, rotations
  - arithmetic, logic
- Vector operations
  - Extract, insert, shuffle
- Memory access & addressing
  - allocation, load, store, fences, getelementptr
- Conversions
  - Zext, sext, ...
- Calling
- Comparisons
  - lcmp, fcmp, ...
- Exceptions
  - Catchpads, landingpads, etc.
- Virtuals (PHI node)
- Intrinsics (gc, padding, etc.)

# Bitcode Instructions

- All instructions are strongly typed and types are almost always explicit in the IR
- Each instruction can be named, in which case the variable bearing the result of the instruction will carry the name

```
%2 = call i32 @min(i32 1,%1)
```

# Bitcode Instructions

- All instructions are strongly typed and types are almost always explicit in the IR
- Each instruction can be named, in which case the variable bearing the result of the instruction will carry the name

```
%2 = call i32 @min(i32 1, %1)
```

Name of the variable holding result of the instruction (if no name is provided, llvm assigns unique number)

# Bitcode Instructions

- All instructions are strongly typed and types are almost always explicit in the IR
- Each instruction can be named, in which case the variable bearing the result of the instruction will carry the name

Llvm instruction (function call)

```
%2 = call i32 @min(i32 1, %1)
```



# Bitcode Instructions

- All instructions are strongly typed and types are almost always explicit in the IR
- Each instruction can be named, in which case the variable bearing the result of the instruction will carry the name

```
%2 = call i32 @min(i32 1, %1)
```

Type of the result of the instruction (explicit)

# Bitcode Instructions

- All instructions are strongly typed and types are almost always explicit in the IR
- Each instruction can be named, in which case the variable bearing the result of the instruction will carry the name

```
%2 = call i32 @min(i32 1, %1)
```

Type of the literal  
argument

# Bitcode Instructions

- All instructions are strongly typed and types are almost always explicit in the IR
- Each instruction can be named, in which case the variable bearing the result of the instruction will carry the name

Value of argument 1 (constant 1)

```
%2 = call i32 @min(i32 1, %1)
```

Value of argument 2 (variable 1)

# Terminators

- `ret`

- return from a function, may or may not return a value

- `ret i32 %3`

- `br`

- Unconditional

- `br label %4`

- Conditional (branches to two basic blocks based on condition)

- `br i1 %2, label %3, label %4`

- Also indirect branch, switch, exceptions throwing & catching



# Binary Operators

- `add, sub, mul`

- Do not care about signedness of operands

- `%1 = add i32 %a, %b`

- `udiv, sdiv`

- Unsigned and signed division of integers

- `fadd, fsub, fmul, fdiv`

- Floating point arithmetics

- `%3 = fadd double %a, %b`

# Memory

- **alloca**

- allocates space for given type on stack and returns a pointer to it

```
%1 = alloca i32, align 4
```

- **load**

- Loads contents of given pointer to register

```
%2 = load i32 i32* %1, align 4
```

- **store**

- stores to a pointer

```
store i32 %I, i32* %1, align 4
```

# Memory

## ➤ `getelementptr`

- address calculation for subelements of aggregate types (array, structure)
- Different semantics for arrays and structs
- Very important and often misunderstood instruction

```
%A = type { i32, double }
```

```
# type of %1 is %A*
```

```
# A*[3].double becomes
```

```
%3 = getelementptr %A, %A* %1, i32 3
```

```
%4 = getelementptr %A, %A* %1, i32 0, i32 1
```

# Memory

## ➤ getelementptr

- address calculation for subelements of aggregate types (array, structure)

- Different semantics for arrays and structures

For arrays, index of element we want to access

Pointer to the aggregate type

```
%A = type { double }  
# type of %1 is double  
# A*[3].double becomes  
%3 = getelementptr %A, %A* %1, i32 3  
%4 = getelementptr %A, %A* %1, i32 0, i32 1
```

Type we are operating on



# Memory

## ➤ getelementptr

- address calculation for subelements of aggregate types (array, structure)
- Different semantics for arrays and structures
- Very important and often misunderstood

For structs, index of element in declaration

Pointer to the aggregate type

```
}  
# type ...  
# A*[3].double ...  
%3 = getelementptr %A, %A* %1, i32 3  
%4 = getelementptr %A, %A* %1, i32 0, i32 1
```

Type we are operating on

# Memory

## ➤ `getelementptr`

- address calculation for subelements of aggregate types (array, structure)
- Different semantics for arrays and structures
- Very important and often misunderstood instruction

For structs, index of element in declaration

Pointer to the aggregate type

```
}  
# type ...  
# A*[3].double ...  
%3 = getelementptr %A, %A* %1, i32 3  
%4 = getelementptr %A, %A* %1, i32 0, i32 1
```

???

Type we are operating on

# Bitcode Representations


- Traditionally LLVM differentiates between three bitcode representations
  - Human readable LLVM (this is what we have seen so far)  
`%1 = add i32 %a, %b`
  - Binary bitcode (this is what LLVM toolchain members usually pass)  
`0x0 f2 cd 0a 0b 54 5a 35 12 7e 2f`
  - C++ API (this is what frontend developers use to construct the IR)  

```
auto b = llvm::BasicBlock::Create(llvm::GetGlobalContext(), "", f);  
auto li = new llvm::LoadInst(ptr, "", false, b);  
llvm::ReturnInst::Create(llvm::getGlobalContext(), li, b);
```



# C++ API

(creating bitcode programmatically)



# Everything is a class, something is a pointer

- Modules, Functions and BasicBlocks
  - `llvm::Module`, `llvm::Function`, `llvm::BasicBlock`
- Types, values and instructions are represented by classes
  - `llvm::Type`, `llvm::Value` and `llvm::Instruction`
- Usually pointers to the classes are expected
- Many (but not all) of the classes should not be created using constructors, but provide static `Create` methods

# Types

- Each type must be specified and have corresponding class
- Types are common to all modules in llvm

```
auto tv=llvm::Type::getVoidTy(llvm::getGlobalContext())
auto ti=llvm::IntegerType::get(context, 32)
auto td=llvm::Type::getDoubleTy(context)
auto ts=llvm::StructType::create(context, "name")
std::vector<llvm::Type *> fields = { ti, ti, td }
ts->setBody(fields, false);
```

# Types

- Each type must be specified and have corresponding class
- Types are common to all modules in llvm

Often, context is required (in case of multiple llvm instances, we can always use global context)

```
auto tv=llvm::Type::getVoidTy(llvm::getGlobalContext())
auto ti=llvm::IntegerType::get(context, 32)
auto td=llvm::Type::getDoubleTy(context)
auto ts=llvm::StructType::create(context, "name")
std::vector<llvm::Type *> fields = { ti, ti, td }
ts->setBody(fields, false);
```

# Types

- Each type must be specified and have corresponding class
- Types are common to all modules in llvm

C++ void

```
auto tv=llvm::Type::getVoidTy(llvm::getGlobalContext())
auto ti=llvm::IntegerType::get(context, 32)
auto td=llvm::Type::getDoubleTy(context)
auto ts=llvm::StructType::create(context, "name")
std::vector<llvm::Type *> fields = { ti, ti, td }
ts->setBody(fields, false);
```



# Types

- Each type must be specified and have corresponding class
- Types are common to all modules in llvm

C++ int or unsigned

```
auto tv=llvm::Type::getVoidTy(llvm::getGlobalContext())
auto ti=llvm::IntegerType::get(context, 32)
auto td=llvm::Type::getDoubleTy(context)
auto ts=llvm::StructType::create(context, "name")
std::vector<llvm::Type *> fields = { ti, ti, td }
ts->setBody(fields, false);
```

# Types

- Each type must be specified and have corresponding class
- Types are common to all modules in llvm

C++ double

```
    tv=llvm::Type::getVoidTy(llvm::getGlobalContext())  
    auto ti=llvm::IntegerType::get(context, 32)  
    auto td=llvm::Type::getDoubleTy(context)  
    auto ts=llvm::StructType::create(context, "name")  
    std::vector<llvm::Type *> fields = { ti, ti, td }  
    ts->setBody(fields, false);
```

# Types

- Each type must be specified and have corresponding class
- Types are common to all modules in llvm

```
auto tv=llvm::Type::getVoidTy(llvm::getGlobalContext())
auto ti=llvm::IntegerType::get(context, 32)
auto td=llvm::Type::getDoubleTy(context)
auto ts=llvm::StructType::create(context, "name")
std::vector<llvm::Type *> fields = { ti, ti, td }
ts->setBody(fields, false);
```

```
struct { int a; int b;
        double c; }
```

Packed?



# Creating constants

- Constants inherit from `llvm::Value` and can be used as arguments to instructions
- As types, constants live outside modules

```
llvm::ConstantInt::get(context, llvm::APInt(32, 1,  
false));
```

```
llvm::ConstantFP::get(context, llvm::APFloat(3.14))
```

# Creating constants

- Constants inherit from `llvm::Value` and can be used as arguments to instructions
- As types, constants live outside modules

```
llvm::ConstantInt::get(context, llvm::APInt(32, 1,  
false))
```

size

value

signed?

```
llvm::ConstantFP::get(context, llvm::APFloat(3.14))
```



# Creating constants

- Constants inherit from `llvm::Value` and can be used as arguments to instructions
- As types, constants live outside modules

```
llvm::ConstantInt::get(context, llvm::APInt(32, 1,  
false))
```

```
llvm::ConstantFP::get(context, llvm::APFloat(3.14))
```



# Creating a function

```
int min(int i, int j) {  
    return i < j ? i : j;  
}
```



# Creating a function

```
auto m = llvm::Module::Create("name", context);  
auto ft = llvm::FunctionType::get(ti, { ti, ti },  
false);  
auto f = llvm::Function::Create(ft,  
llvm::GlobalValue::ExternalLinkage, "min", m);  
f->setCallingConvention(llvm::CallingConv::C);
```



# Creating a function

First create a module with given name

```
auto m = llvm::Module::Create("name", context);  
auto ft = llvm::FunctionType::get(ti, { ti, ti },  
false);  
auto f = llvm::Function::Create(ft,  
llvm::GlobalValue::ExternalLinkage, "min", m);  
f->setCallingConvention(llvm::CallingConv::C);
```

# Creating a function

The function must have a type, in our case  
`int (*ptr)(int, int)`

`text);`

```
auto ft = llvm::FunctionType::get(ti, { ti, ti },  
false);
```

varargs?

```
auto f = llvm::Function::Create(ft,  
llvm::GlobalValue::ExternalLinkage, "min", m);  
f->setCallingConvention(llvm::CallingConv::C);
```

# Creating a function

```
auto m = llvm::Module::Create("name", context);
```

Create the function object

type::get

type

```
auto f = llvm::Function::Create(ft,  
    llvm::GlobalValue::ExternalLinkage, "min", m);
```

symbol visibility

```
f->setCallingConvention(llvm::CallingConv::C);
```

name of the function  
(unique)

Module the function  
belongs to

# Creating a function

```
auto m = llvm::Module::Create("name", context);  
auto ft = llvm::FunctionType::get(ti, { ti, ti },  
false);  
auto f = llvm::Function::Create(ft,  
llvm::GlobalValue::ExternalLinkage, "min", m);  
f->setCallingConvention(llvm::CallingConv::C);
```

sets C calling convention for the function

# Creating a function

```
auto m = llvm::Module::Create("name", context);  
auto ft = llvm::FunctionType::get(ti, { ti, ti },  
false);  
auto f = llvm::Function::Create(ft,  
llvm::GlobalValue::ExternalLinkage, "min", m);  
f->setCallingConvention(llvm::CallingConv::C);
```

At this point we have created a proper function declaration in LLVM IR. The function can be called in the module, but it does not have any code in it.



# Creating code

```
auto args = f->arg_begin();
llvm::Value * first = args++;
llvm::Value * second = args;

auto b = llvm::BasicBlock::Create(context, "first", f);
auto cmp = new llvm::ICmpInst(*b,
    llvm::ICmpInst::ICMP_SLT, first, second);

auto lt = llvm::BasicBlock::Create(context, "lt", f);
auto gte = llvm::BasicBlock::Create(context, "gte", f);
llvm::BranchInst::Create(lt, gte, cmp, b);

llvm::ReturnInst::Create(context, lt, first);

llvm::ReturnInst::Create(context, gte, second);
```

# Creating code

Create shorthand values  
for function arguments

```
auto args = f->arg_begin();  
llvm::Value * first = args++;  
llvm::Value * second = args;  
  
auto b = llvm::BasicBlock::Create(context, "first", f);  
auto cmp = new llvm::ICmpInst(*b,  
    llvm::ICmpInst::ICMP_SLT, first, second);  
  
auto lt = llvm::BasicBlock::Create(context, "lt", f);  
auto gte = llvm::BasicBlock::Create(context, "gte", f);  
llvm::BranchInst::Create(lt, gte, cmp, b);  
  
llvm::ReturnInst::Create(context, lt, first);  
  
llvm::ReturnInst::Create(context, gte, second);
```

# Creating code

Create the first  
basic block

```
auto args = f->arg_begin();  
llvm::Value * first = args++;  
llvm::Value * second = args;  
  
auto b = llvm::BasicBlock::Create(context, "first", f);  
auto cmp = new llvm::ICmpInst(*b,  
    llvm::ICmpInst::ICMP_SLT, first, second);  
  
auto lt = llvm::BasicBlock::Create(context, "lt", f);  
auto gte = llvm::BasicBlock::Create(context, "gte", f);  
llvm::BranchInst::Create(cmp, lt, gte, b);  
  
llvm::ReturnInst::Create(context, lt, first);  
  
llvm::ReturnInst::Create(context, gte, second);
```

Insert signed less than comparison of the  
arguments at the end of the basic block



# Creating code

```
auto args = f->arg_begin();  
llvm::Value * first = args++;
```

Create basic blocks for  $i < j$  and  $i \geq j$  cases

```
llvm::ICmpInst * cmp = llvm::ICmpInst::Create(llvm::CmpInst::Cmp_SLE, first, second),  
auto lt = llvm::BasicBlock::Create(context, "lt", f);  
auto gte = llvm::BasicBlock::Create(context, "gte", f);  
llvm::BranchInst::Create(lt, gte, cmp, b);
```

```
llvm::ReturnInst
```

```
llvm::ReturnInst
```

Conditional branch based on the result of the comparison to either lt, or gte basic blocks

# Creating code

```
auto args = f->arg_begin();
llvm::Value * first = args++;
llvm::Value * second = args++;

auto b = llvm::BasicBlock::Create(context, "first", f);
auto cmp = llvm::ICmpInst::Create(llvm::CmpInst::ICMP_EQ, first, second, b);

auto lt = llvm::BasicBlock::Create(context, "lt", f);
auto gte = llvm::BasicBlock::Create(context, "gte", f);
llvm::BranchInst::Create(lt, gte, cmp, b);

llvm::ReturnInst::Create(context, lt, first);

llvm::ReturnInst::Create(context, gte, second);
```

Return the respective values  
in the branches




# Creating code

```
define i32 @min(i32 %1, i32 %2) {  
    %3 = icmp slt i32 %1, %2  
    br i1 %3, label %4, label %5  
  
; <label>:4  
    ret i32 %1  
  
; <label>:5  
    ret i32 %2  
}
```

# Working with PHI nodes

- In the min function, we can create only one exit point from the function
- But it would require us to use a PHI node as we would return either first, or second argument in the last basic block based on where we arrived from



# Working with PHI nodes

```
define i32 @min(i32 %1, i32 %2) {  
    %3 = icmp slt i32 %1, %2  
    br i1 %3, label %4, label %5  
  
; <label>:4  
    br label %6  
  
; <label>:5  
    br label %6  
  
; <label>:6  
    %7 = phi i32 [ %1 %4 ], [ %2 %5 ]  
    ret i32 %7  
}
```

# Working with PHI nodes

```
define i32 @min(i32 %1, i32 %2) {  
    %3 = icmp slt i32 %1, %2  
    br i1 %3, label %4, label %5  
  
; <label>:4  
    br label %6  
; <label>:5  
    br label %6  
; <label>:6  
    %7 = phi i32 [ %1 %4 ], [ %2 %5 ]  
    ret i32 %7  
}
```

If arriving from block  
%4, value will be %1

If arriving from block  
%5, value will be %2

# Working with PHI nodes

```
auto args = f->arg_begin();
llvm::Value * first = args++;
llvm::Value * second = args;

auto b = llvm::BasicBlock::Create(context, "first", f);
auto cmp = new llvm::ICmpInst(*b,
llvm::ICmpInst::ICMP_SLT, first, second);

auto lt = llvm::BasicBlock::Create(context, "lt", f);
auto gte = llvm::BasicBlock::Create(context, "gte", f);
auto end = llvm::BasicBlock::Create(context, "end", f);

llvm::BranchInst::Create(lt, gte, cmp, b);
llvm::BranchInst::Create(end, lt);
llvm::BranchInst::Create(end, gte);

auto phi = llvm::PHINode::Create(ti, 2, "", end);
phi->addIncoming(first, lt);
phi->addIncoming(second, gte);

llvm::ReturnInst::Create(context, end, phi);
```

# Working with PHI nodes

```
auto args = f->arg begin();
```

Create the ending basic block and jumps from lt and gte blocks to it (they will be empty)

```
auto lt = llvm::BasicBlock::Create(context, "lt", f);  
auto gte = llvm::BasicBlock::Create(context, "gte", f);  
auto end = llvm::BasicBlock::Create(context, "end", f);
```

```
llvm::BranchInst::Create(lt, gte, cmp, b);  
llvm::BranchInst::Create(end, lt);  
llvm::BranchInst::Create(end, gte);
```

```
auto phi = llvm::PHINode::Create(ti, 2, "", end);  
phi->addIncomming(first, lt);  
phi->addIncomming(second, gte);
```

```
llvm::ReturnInst::Create(context, end, phi);
```



# Working with PHI nodes

```
auto args = f->arg_begin();
llvm::Value * first = args++;
llvm::Value * second = args;

auto b = llvm::BasicBlock::Create(context, "first", f);
auto cmp = new llvm::ICmpInst(*b,
llvm::ICmpInst::ICMP_SLT, first, second);

auto lt = llvm::BasicBlock::Create(context, "lt", f);
auto gte = llvm::BasicBlock::Create(context, "gte", f);
auto end = llvm::BasicBlock::Create(context, "end", f);

// BranchInst::Create(branches, cmp, b);
// BranchInst::Create(branches, gte);

auto phi = llvm::PHINode::Create(ti, 2, "", end);
phi->addIncoming(first, lt);
phi->addIncoming(second, gte);

llvm::ReturnInst::Create(context, end, phi);
```

Create phi node in the last basic block,  
reserve 2 incoming edges

Add incoming  
edges



# Command Line

(playing with LLVM)



# Exporting bitcode

```
clang++ -S -emit-llvm test.cpp -o test.ll
```

- ▶ emits bitcode in human readable form, created by clang frontend from given file
- ▶ Slightly more complex than the examples we have seen so far
  - ▶ Function attributes
  - ▶ Metadata
  - ▶ Beware of C++ constructs, name mangling, etc.



# Optimizing bitcode

```
opt test.ll -o test.opt.ll -S -mem2reg -dce -constprop
```

- ▀ Invokes llvm optimizer on given bitcode file, runs specified passes and outputs the resulting bitcode

```
opt --help
```

- ▀ Shows all passes opt understands



# C++ API help

```
llc -march=cpp test.ll -o test.ll.cpp
```

- ▶ Takes bitcode in test.ll and compiles it using specified target
- ▶ The cpp target is very useful as its output is the C++ API calls required to create the input bitcode
- ▶ Definitely more complex than what we have seen so far, but fairly human readable for small examples



Q & A

Thank you