

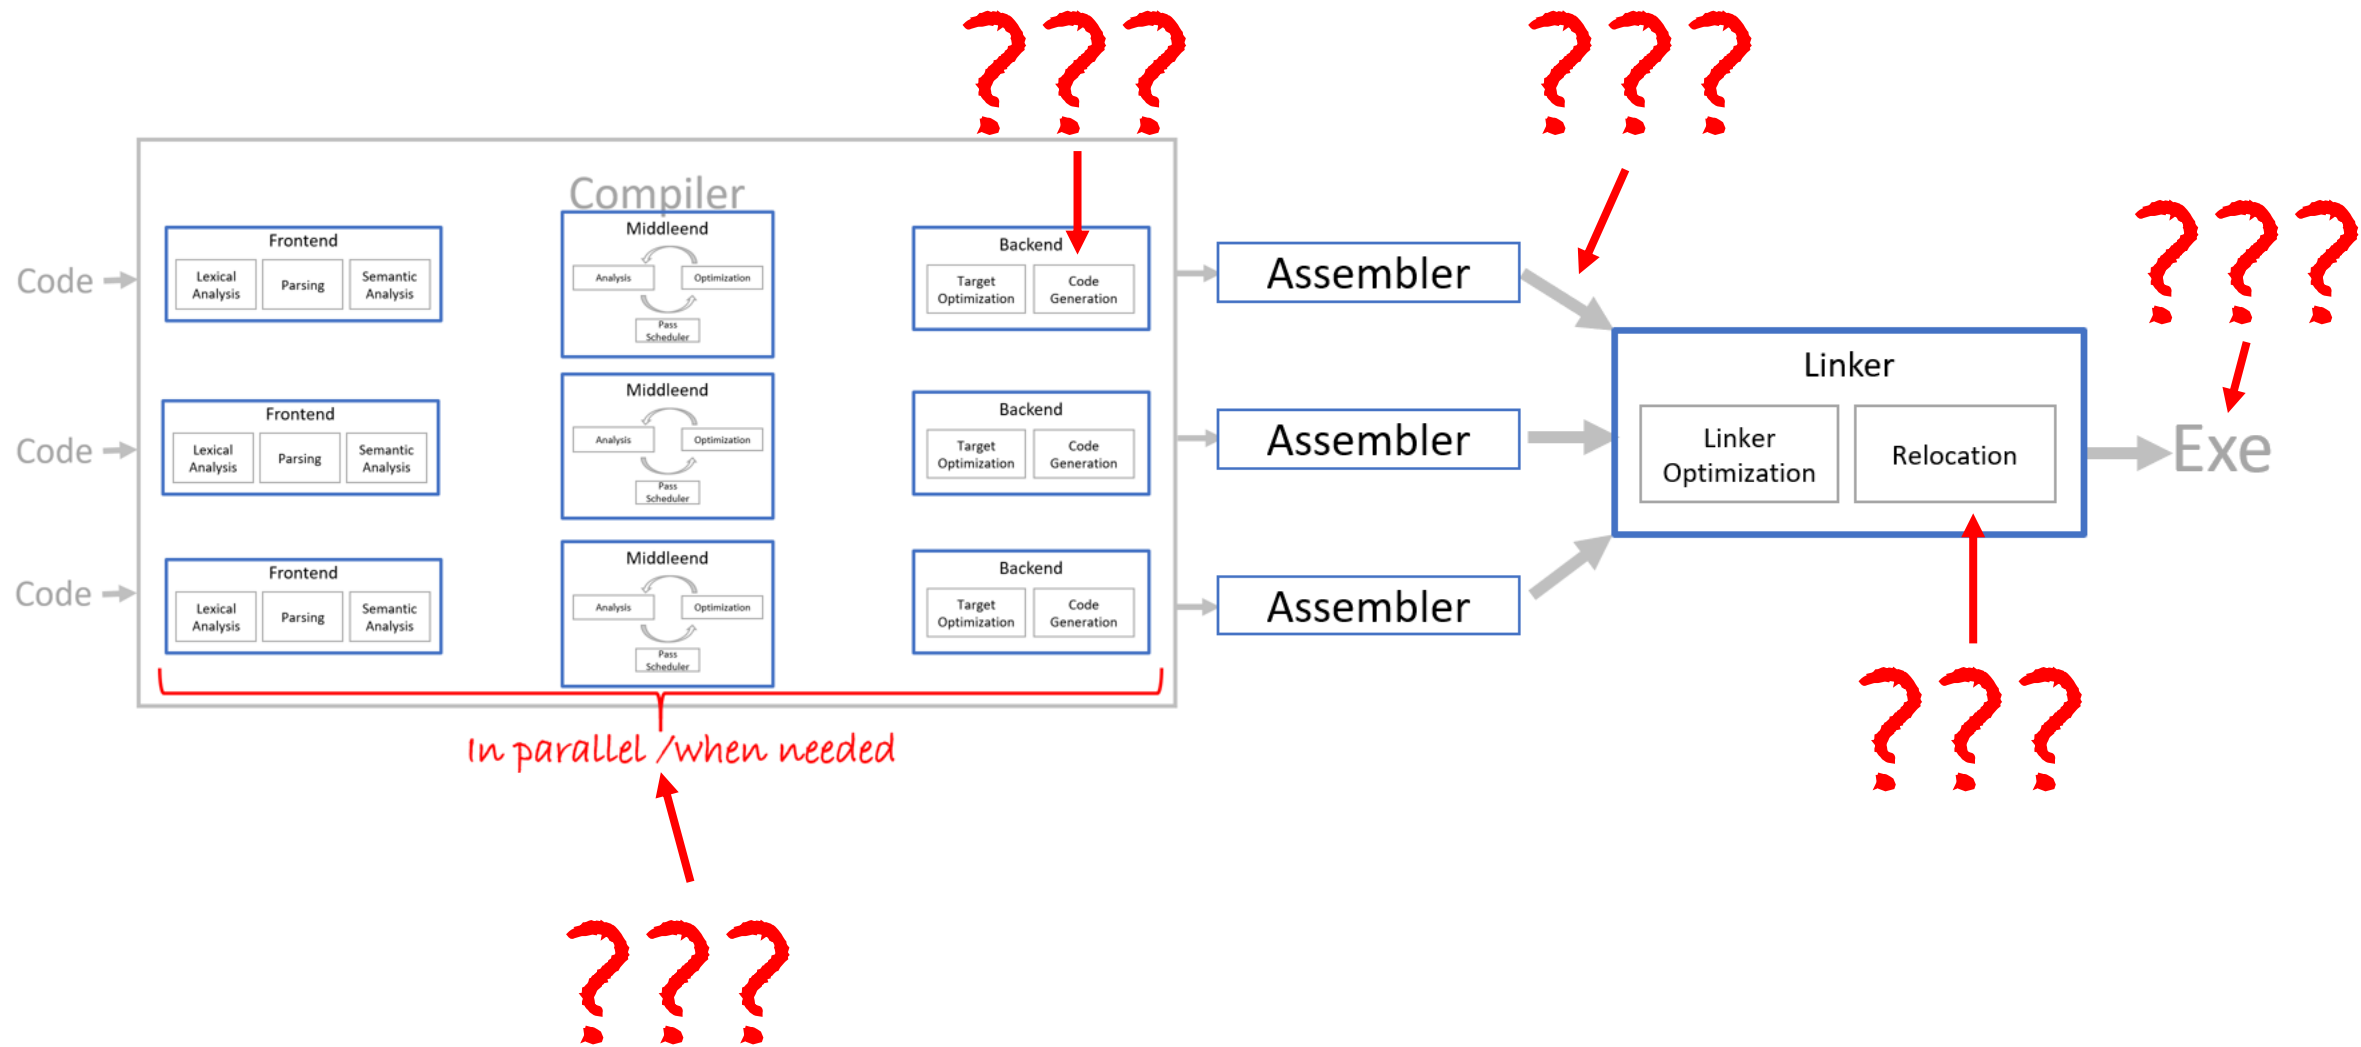
Principles of Compiled Code

NI(E)-GEN, Spring 2021

<https://courses.fit.cvut.cz/NI-GEN>



So we have a compiler...



Executables

Executable

- contains binary data of various kinds partitioned in sections (.text, .data, .rodata, etc.)

.text:

mov eax, hello_world

...

.rodata

hello_world:

“hello all, this is a text”, 0

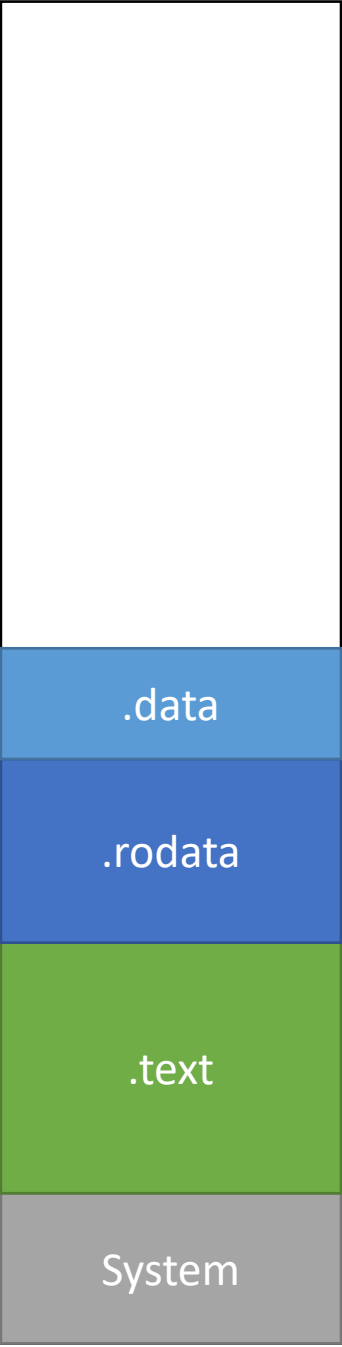
.data

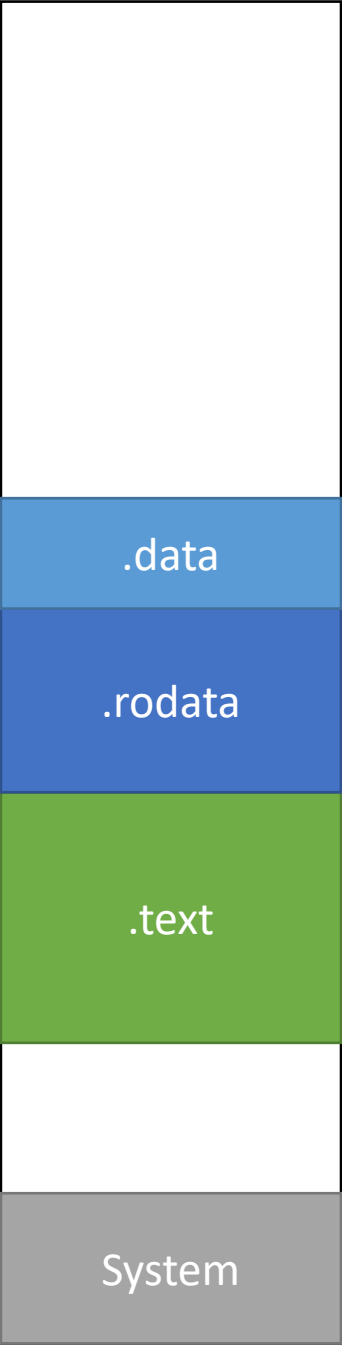
global-var:

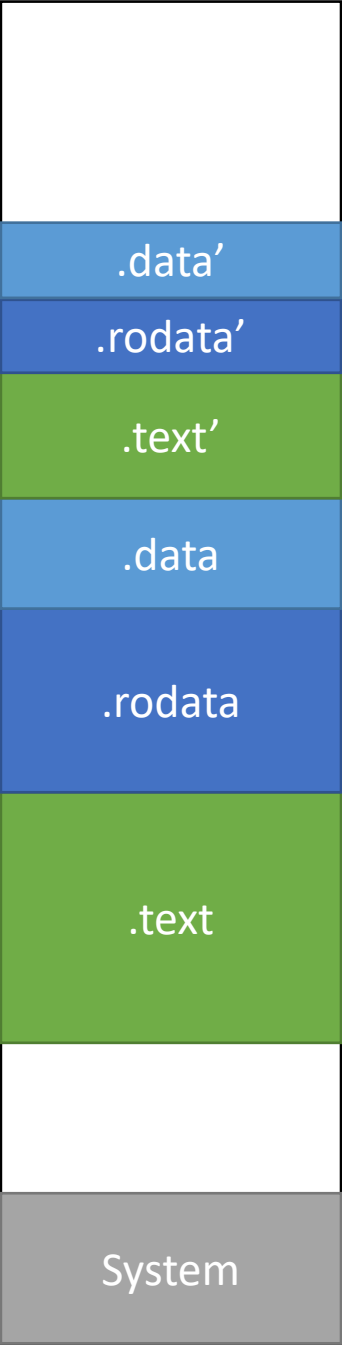
0x67

Executable

- contains binary data of various kinds partitioned in sections (.text, .data, .rodata, etc.)
- when executed, the section contents are copied to memory








Relocation

- absolute addresses of symbols cannot be known statically
- executable contains a relocation table
 - address to be patched
 - target symbol
 - patch type
- the loader then updates the section contents based on the relocation table once the section starts are known

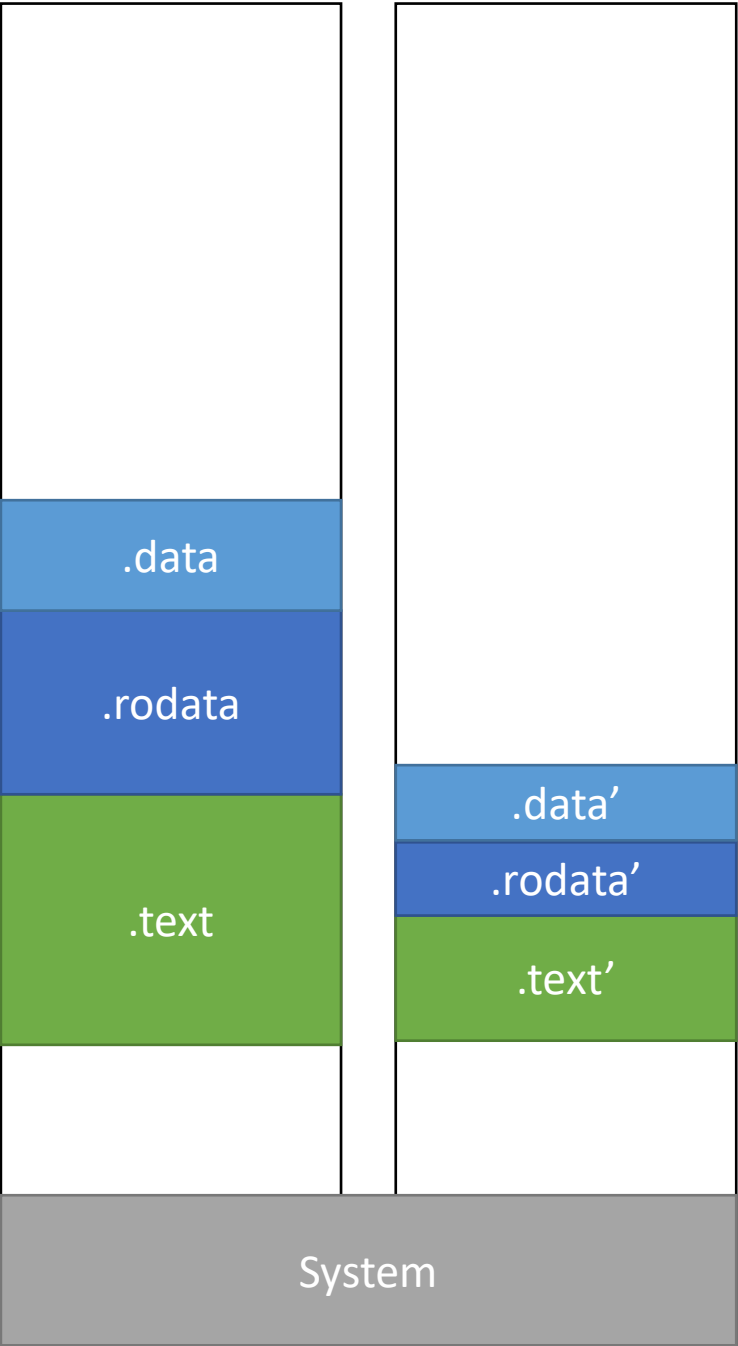
Executable

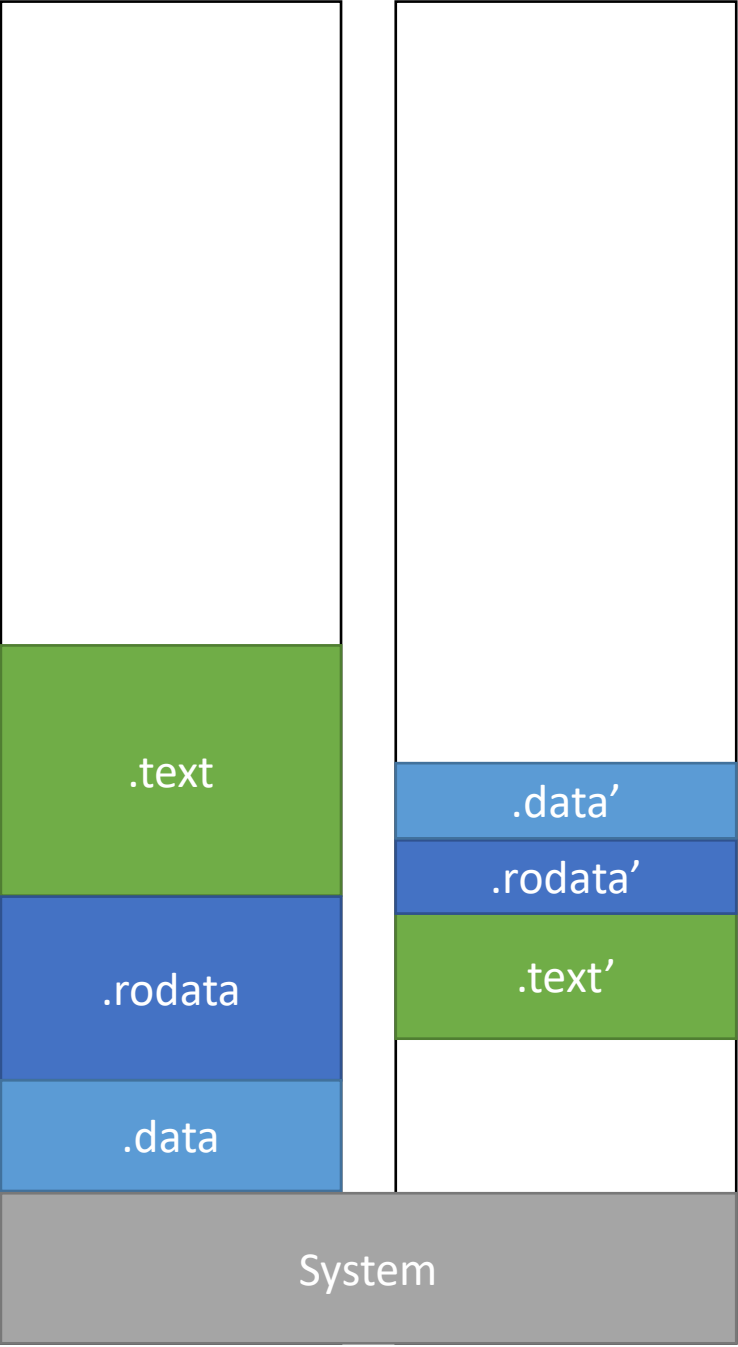
- contains binary data of various kinds partitioned in sections (.text, .data, .rodata, etc.)
- when executed, the section contents are copied to memory
- and relocated



Executable

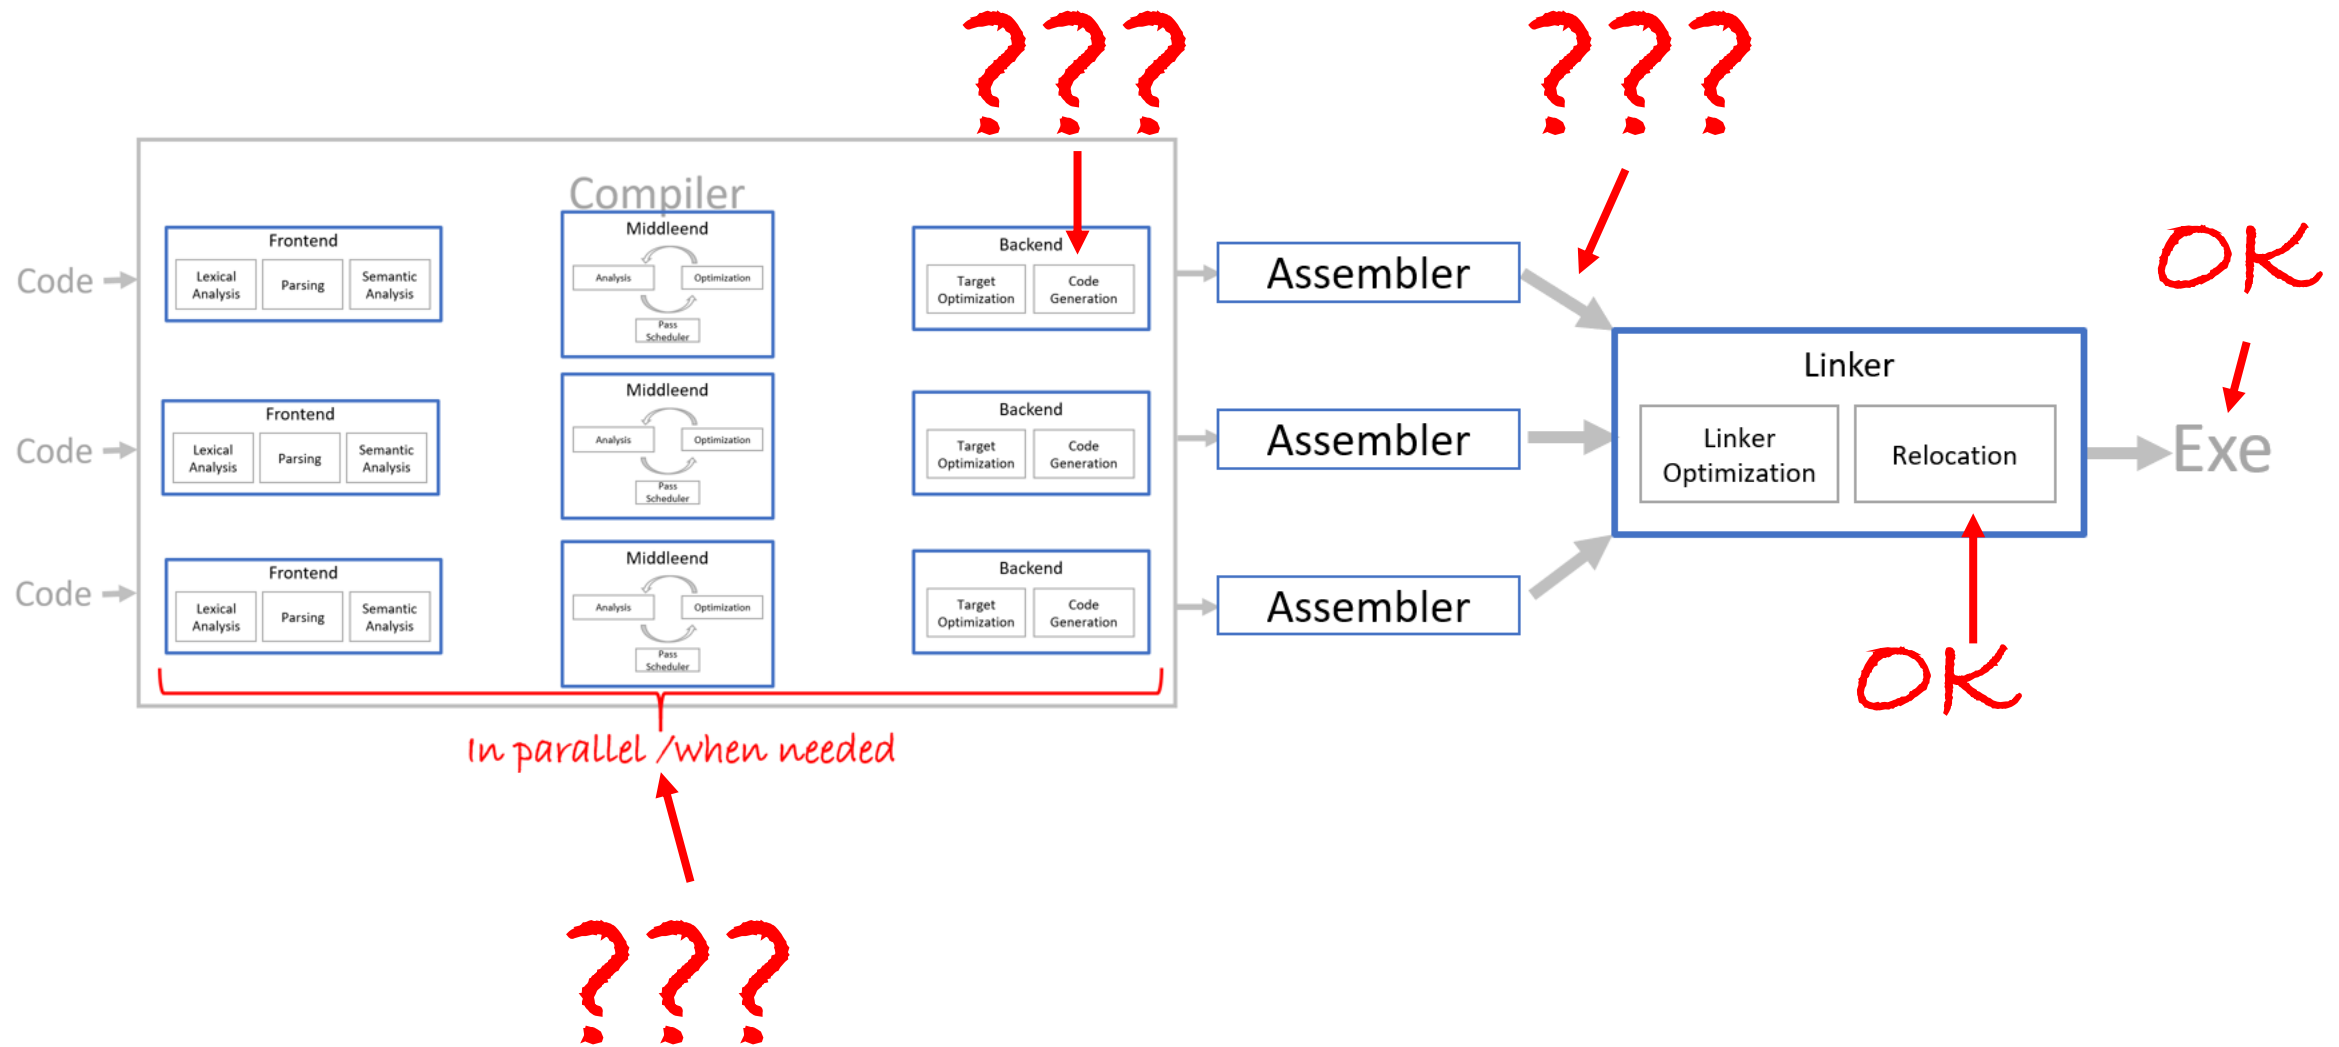
- contains binary data of various kinds partitioned in sections (.text, .data, .rodata, etc.)
- when executed, the section contents are copied to memory
- and relocated (if necessary, such as libraries)





Executable

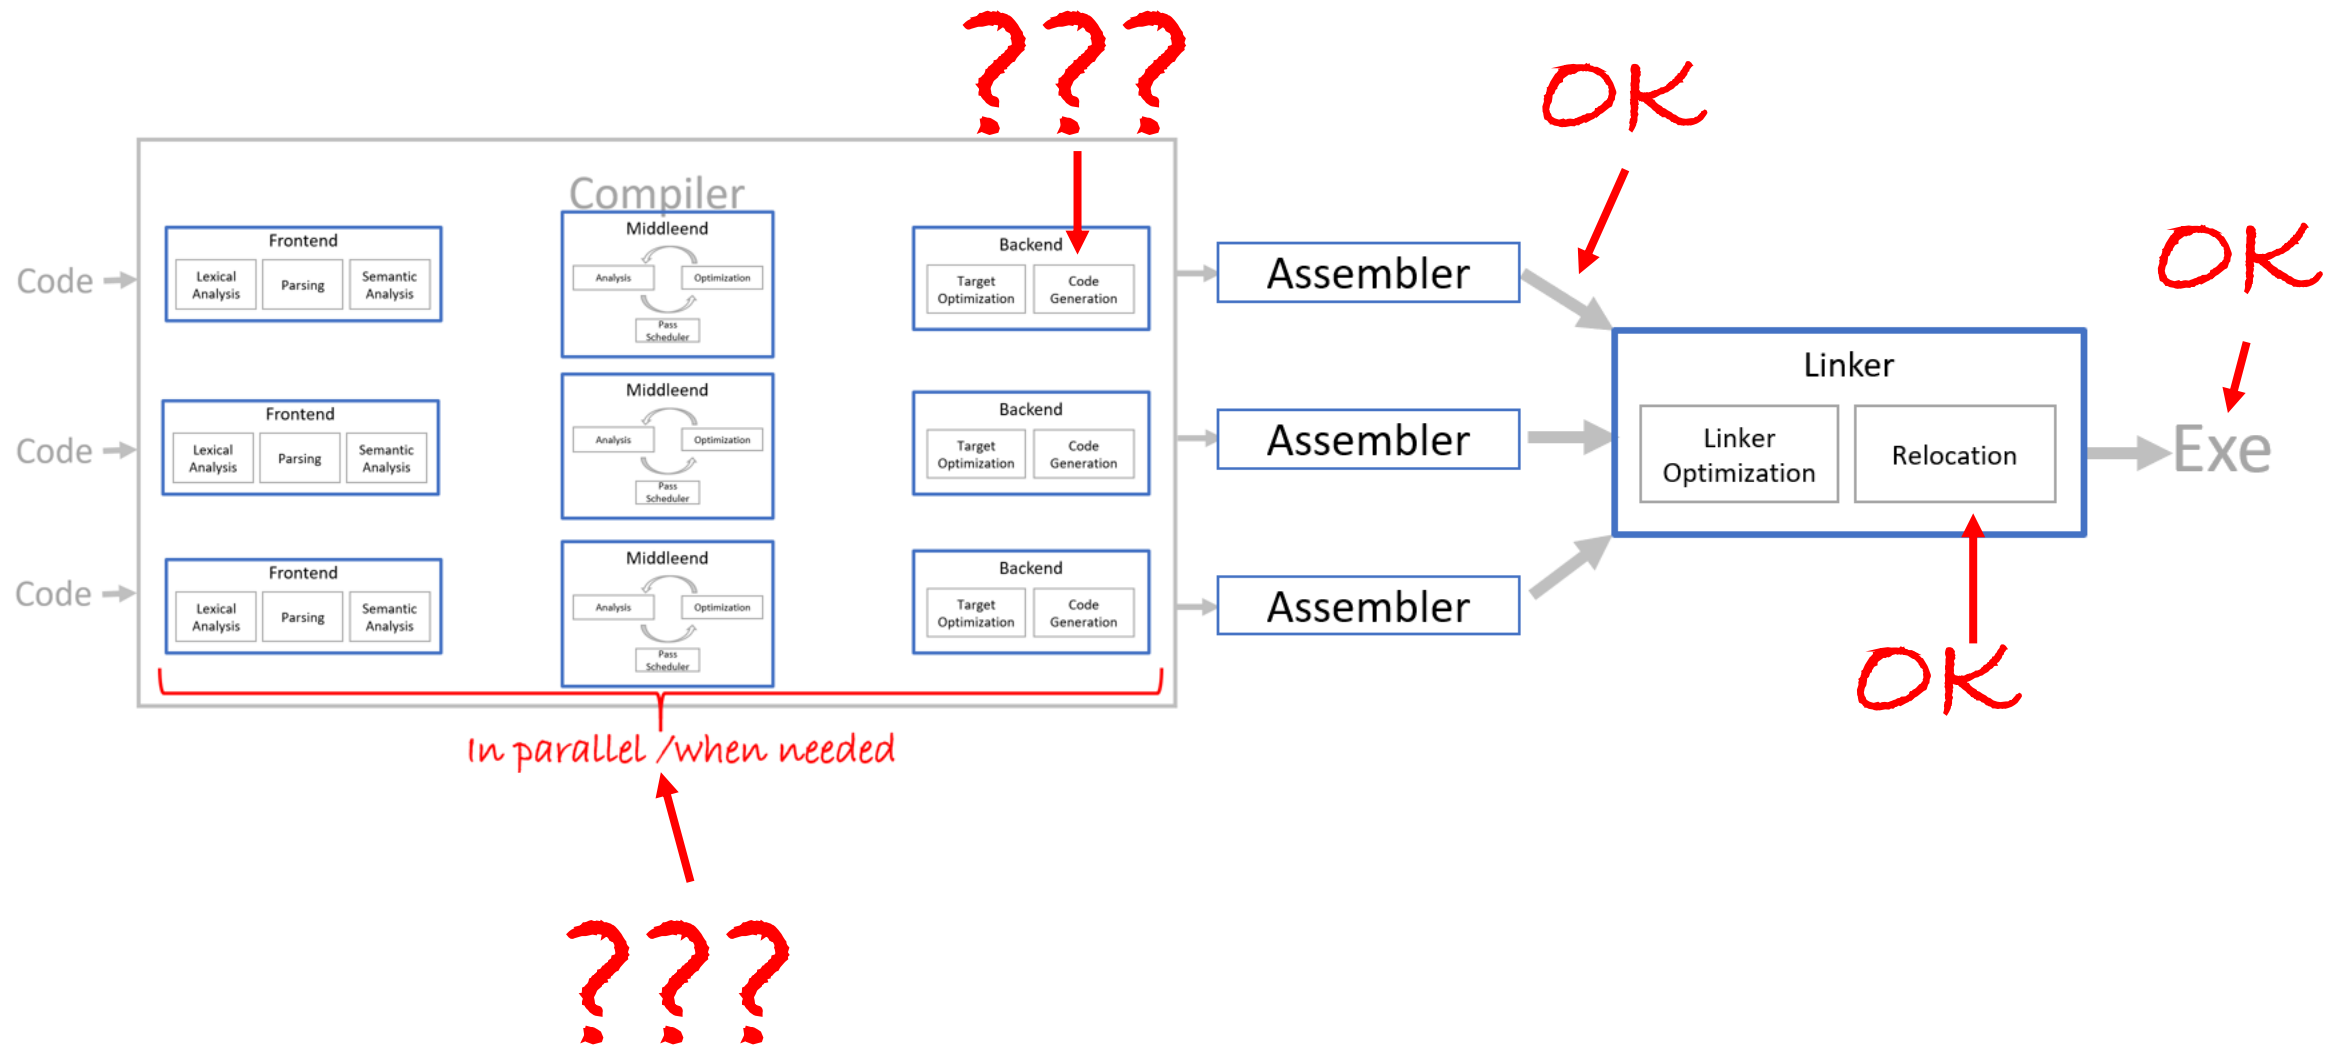
- contains binary data of various kinds partitioned in sections (.text, .data, .rodata, etc.)
- when executed, the section contents are copied to memory
- and relocated (if necessary, such as libraries, or ASLR)



Object Files

Object Files

- like executable
- + more sections, - loader
- a lot more relocations
- more information so that link-time optimization can be performed



Source Granularity

Source Granularity

- project
- executable & libraries
- files
- functions
- basic blocks
- statements

Basic Block

- code sequence with single entry and single leave point (terminating instruction)
- once first instruction in a basic block gets executed, all instructions in basic block will execute sequentially

↖ no control flow


```
if (a < 0) {  
    a = a + 1;  
} else {  
    b = b + 1;  
}  
c = a + b;
```

```
if (a < 0) {  
    a = a + 1;  
} else {  
    b = b + 1;  
}  
c = a + b;
```

```
B0: cmp ax, 0 // a in ax  
    jge B2  
B1: add ax, 1  
    jmp B3  
B2: add bx, 1 // b in bx  
B3: mov cx, ax // c in cx  
    add cx, bx
```

```
if (a < 0) {  
    a = a + 1;  
} else {  
    b = b + 1;  
}  
c = a + b;
```

```
B0: cmp ax, 0 // a in ax  
    jge B2  
B1: add ax, 1  
    jmp B3  
B2: add bx, 1 // b in bx  
    jmp B3  
B3: mov cx, ax // c in cx  
    add cx, bx
```

BI-PPA, BI-PJP

Code Generation

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

f_min: # functions are labels

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

```
f_min:  
    cmp ax, bx # arguments in  
    jl x_less # ax and bx regs  
  
x_less:
```

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

```
f_min:  
    cmp ax, bx  
    jl x_less  
    mov ax, bx # result  
    ret # passed in ax reg  
x_less:
```


Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

```
f_min:  
    cmp ax, bx  
    jl x_less  
    mov ax, bx  
    ret  
x_less:  
    ret # already in ax
```

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

```
jmp start # start with start
```

```
f_min:
```

```
    cmp ax, bx
```

```
    jl x_less
```

```
    mov ax, bx
```

```
    ret
```

```
x_less:
```

```
    ret
```

```
start: # start exec here
```

```
    mov ax, 67 # vars in regs
```

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

```
jmp start  
f_min:  
    cmp ax, bx  
    jl x_less  
    mov ax, bx  
    ret  
x_less:  
    ret  
start:  
    mov ax, 67  
    mov bx, 89
```

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

```
jmp start  
f_min:  
    cmp ax, bx  
    jl x_less  
    mov ax, bx  
    ret  
x_less:  
    ret  
start:  
    mov ax, 67  
    mov bx, 89  
    call f_min # luck:) -> ax
```

Code Generation

```
int min(int x, int y) {  
    return x < y ? x : y;  
}
```

```
int a = 67;  
int b = 89;  
print(min(a, b));
```

```
jmp start  
f_min:  
    cmp ax, bx  
    jl x_less  
    mov ax, bx  
    ret  
x_less:  
    ret  
start:  
    mov ax, 67  
    mov bx, 89  
    call f_min  
    call print # ax -> ax
```

Code Generation

```
00:    jmp start
      f_min:
04:    cmp ax, bx
08:    jl x_less
0a:    mov ax, bx
0c:    ret
      x_less:
0d:    ret
      start:
0e:    mov ax, 67
12:    mov bx, 89
16:    call f_min
1a:    call print
```

Code Generation

```
00:    jmp 0e
04:    cmp ax, bx
08:    jl 0d
0a:    mov ax, bx
0c:    ret
0d:    ret
0e:    mov ax, 67
12:    mov bx, 89
16:    call 04
1a:    call print
```

← Linker, help!

Memory Layout

AX

BX

CX

DX

PC

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

fib(3)

AX

BX

CX

DX

PC

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

main()

fib(3)

AX

BX

CX

DX

PC

1 argument

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

3 local
variables

recursive

return value

fib(3)

Call Stack

- **Last In – First Out** Structure
- Holds local variables, arguments and return addresses
- Each function's data live in its own area called stack frame
- When function exits, the frame is popped off the stack

AX

BX

CX

DX

PC

SP

BP

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

fib(3)

Stack:



AX

BX

CX

DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC

fib(3)

Stack:

BP

main

start of main()

local vars of main

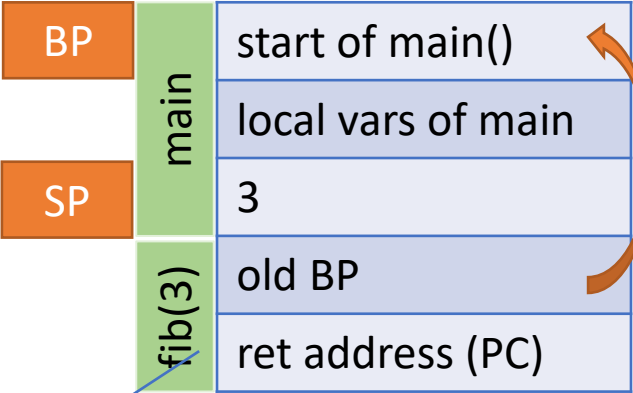
SP

AX
BX
CX
DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC fib(3)

Stack:

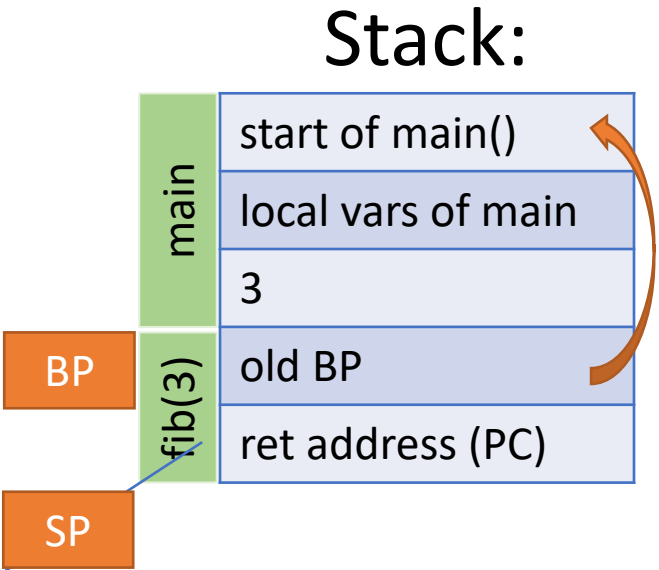


AX
BX
CX
DX

[BP - 1]

PC

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```

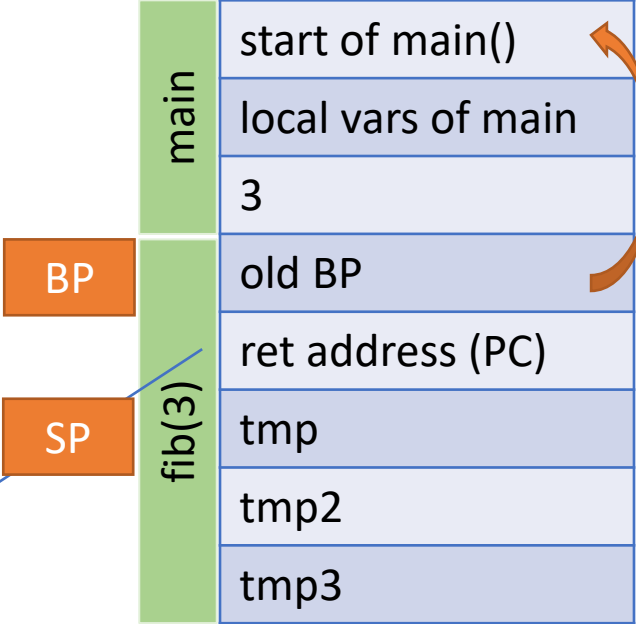


AX
BX
CX
DX

PC `int fib(int n) {
 if (n == 0)
 return 0
 if (n == 1)
 return 1
 int tmp = fib(n - 2)
 int tmp2 = fib(n - 1)
 int tmp3 = tmp1 + tmp2
 return tmp3
}`

fib(3)


Stack:



AX
BX
CX
DX

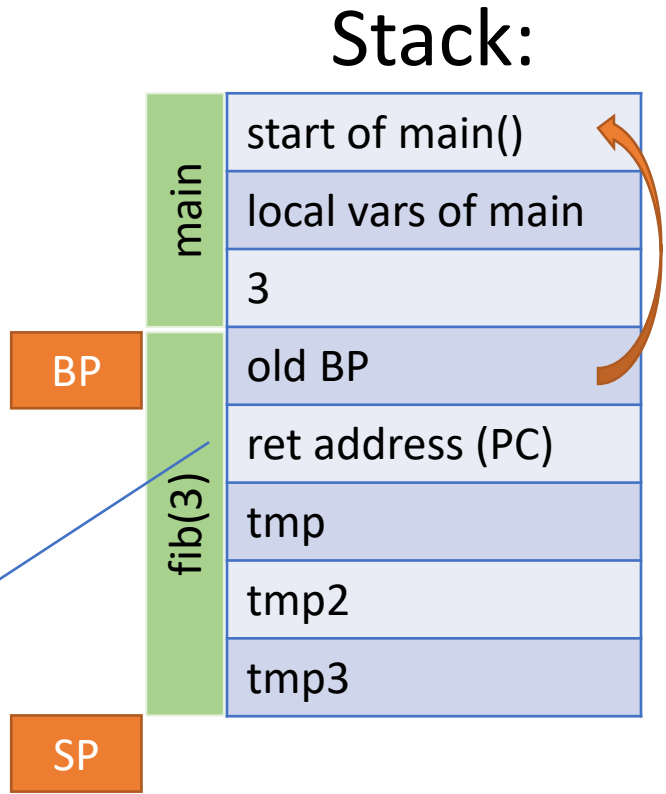
```
PC int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```

[BP + 2]
[BP + 3]
[BP + 4]



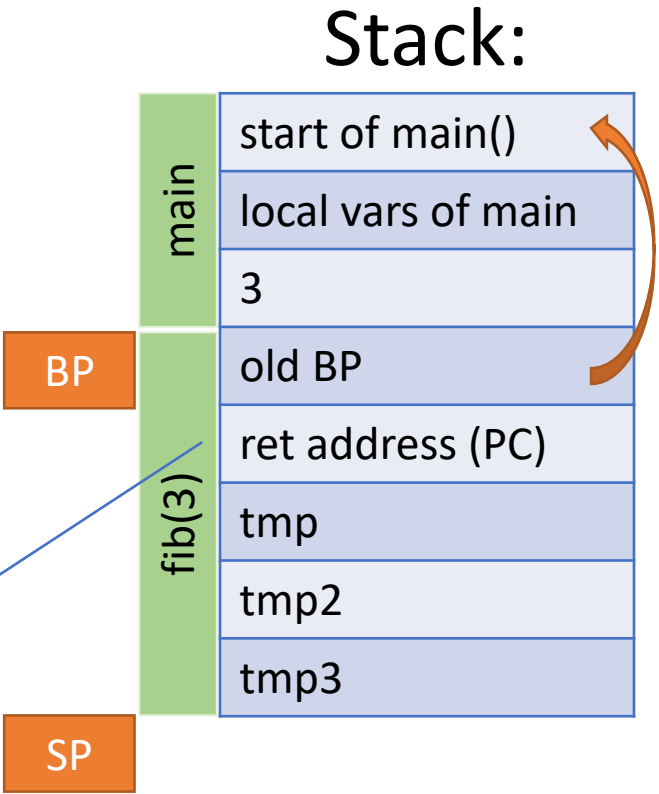
AX
BX
CX
DX

```
int fib(int n) {  
  PC if (n == 0)  
      return 0  
  if (n == 1)  
      return 1  
  int tmp = fib(n - 2)  
  int tmp2 = fib(n - 1)  
  int tmp3 = tmp1 + tmp2  
  return tmp3  
}  
  
fib(3)
```



AX
BX
CX
DX

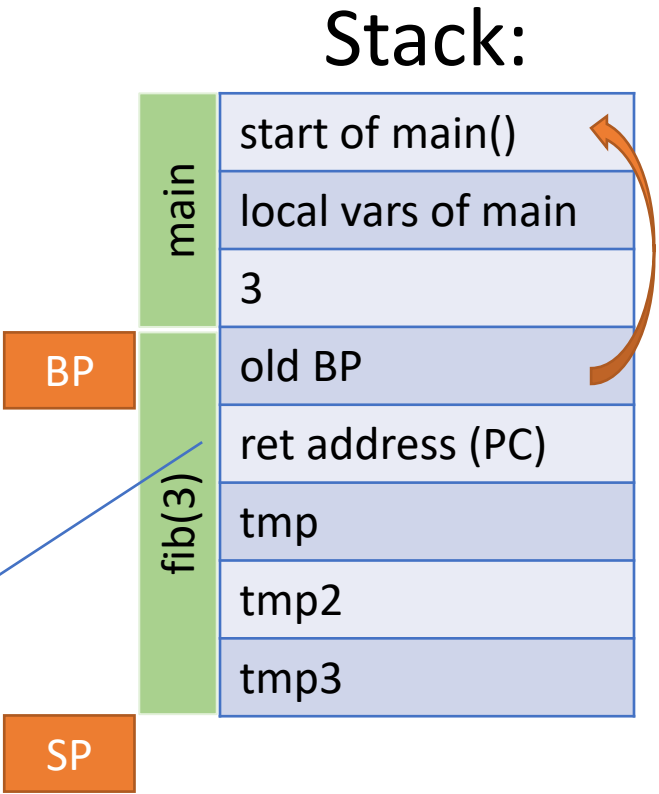
```
int fib(int n) {  
    if (n == 0)  
        return 0  
    PC if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



AX
BX
CX
DX

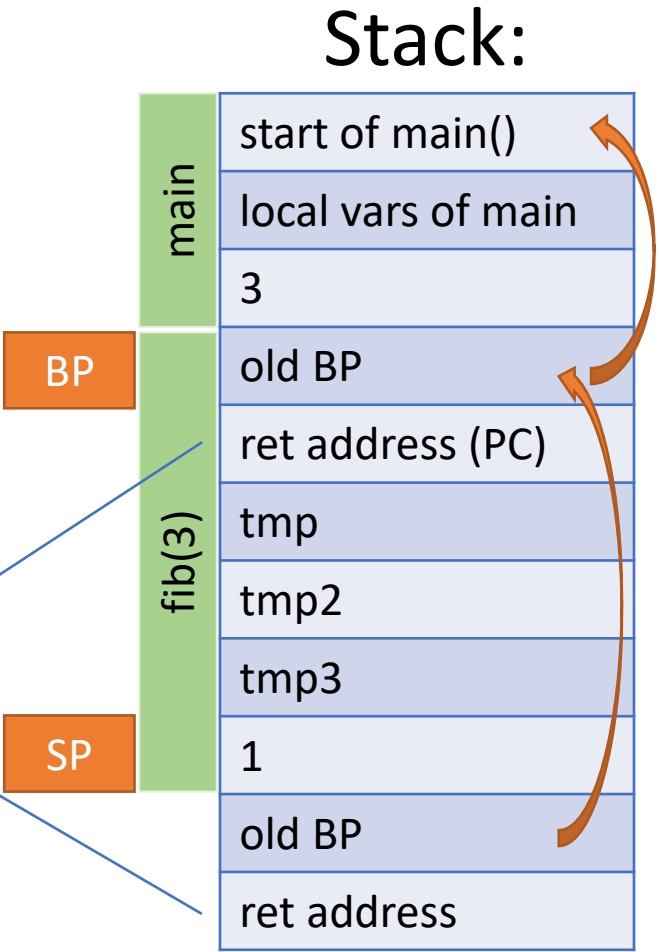
```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

fib(3)



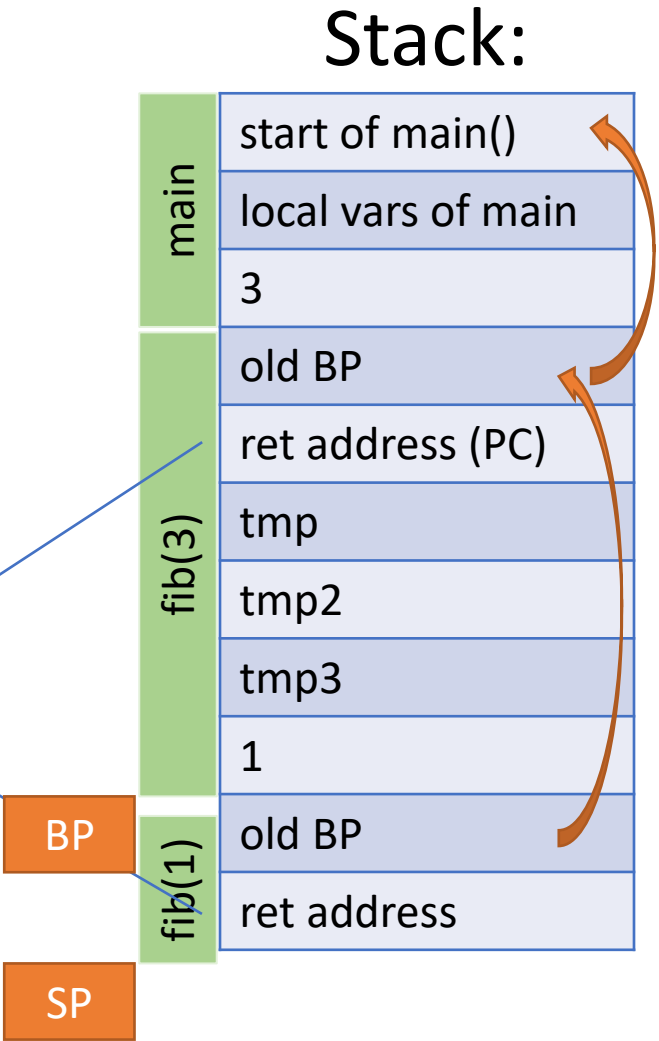
- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    PC int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



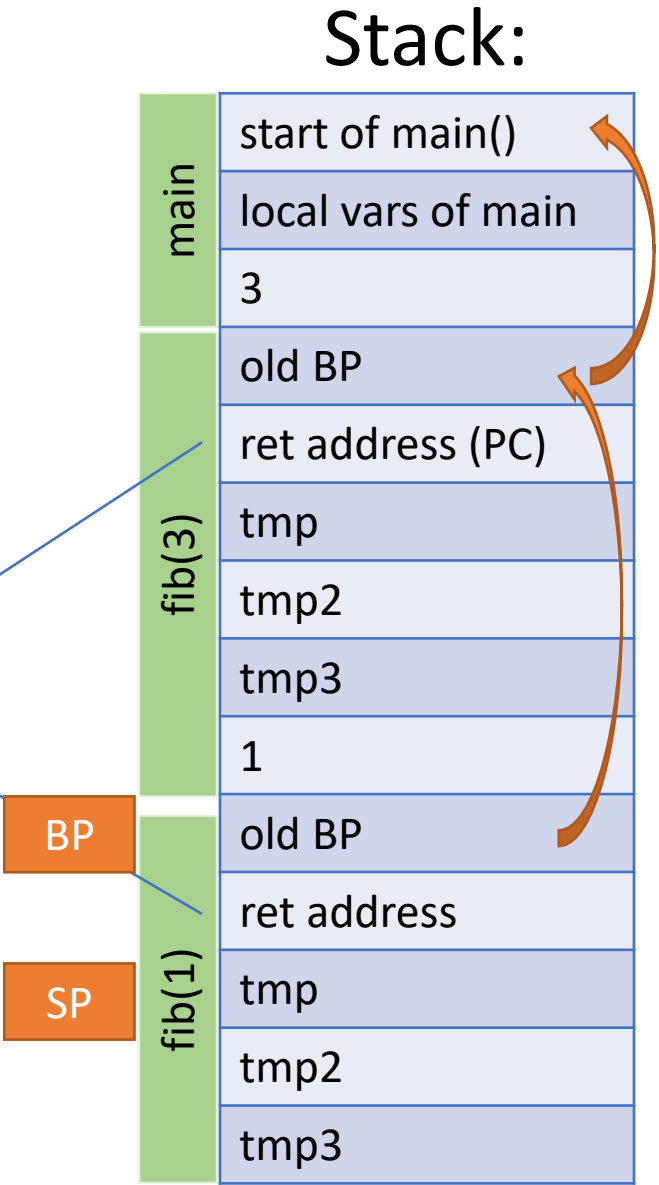
AX
BX
CX
DX

```
PC int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```




- AX
- BX
- CX
- DX

```
PC int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



- AX
- BX
- CX
- DX


```
PC int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



- AX
- BX
- CX
- DX

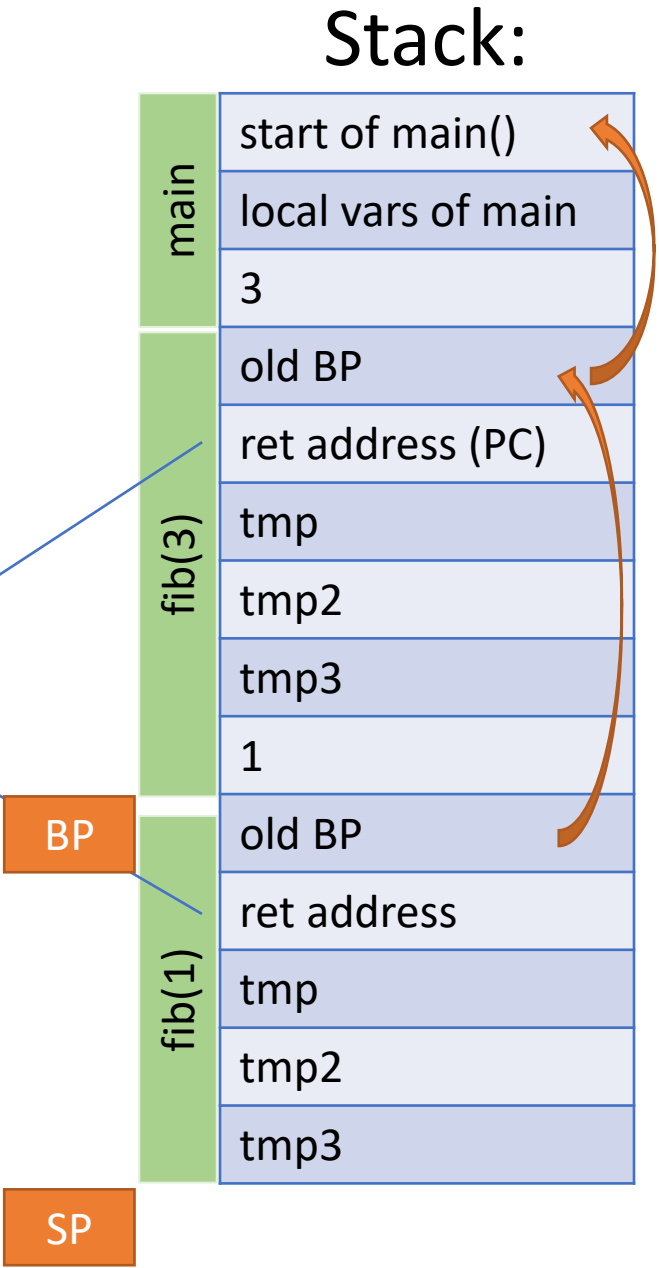
```
int fib(int n) {
  PC if (n == 0)
      return 0
  if (n == 1)
      return 1
  int tmp = fib(n - 2)
  int tmp2 = fib(n - 1)
  int tmp3 = tmp1 + tmp2
  return tmp3
}

fib(3)
```




- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    PC if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



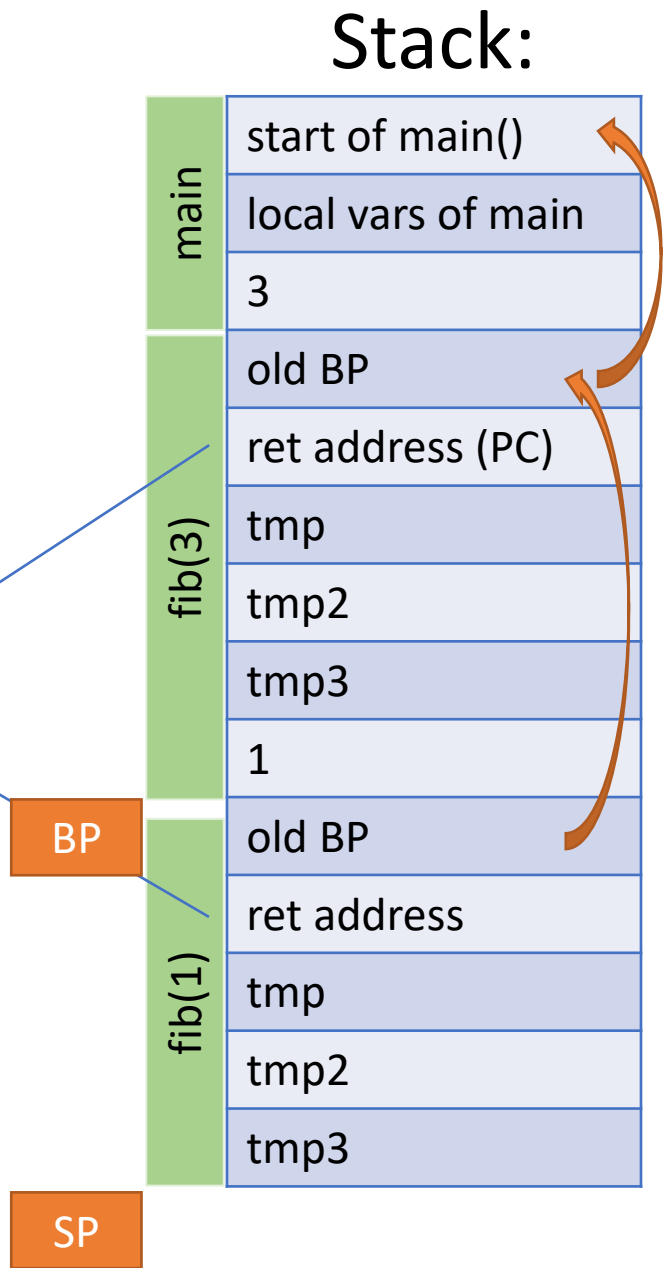
- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        PC return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```





fib(3)



AX

1

BX

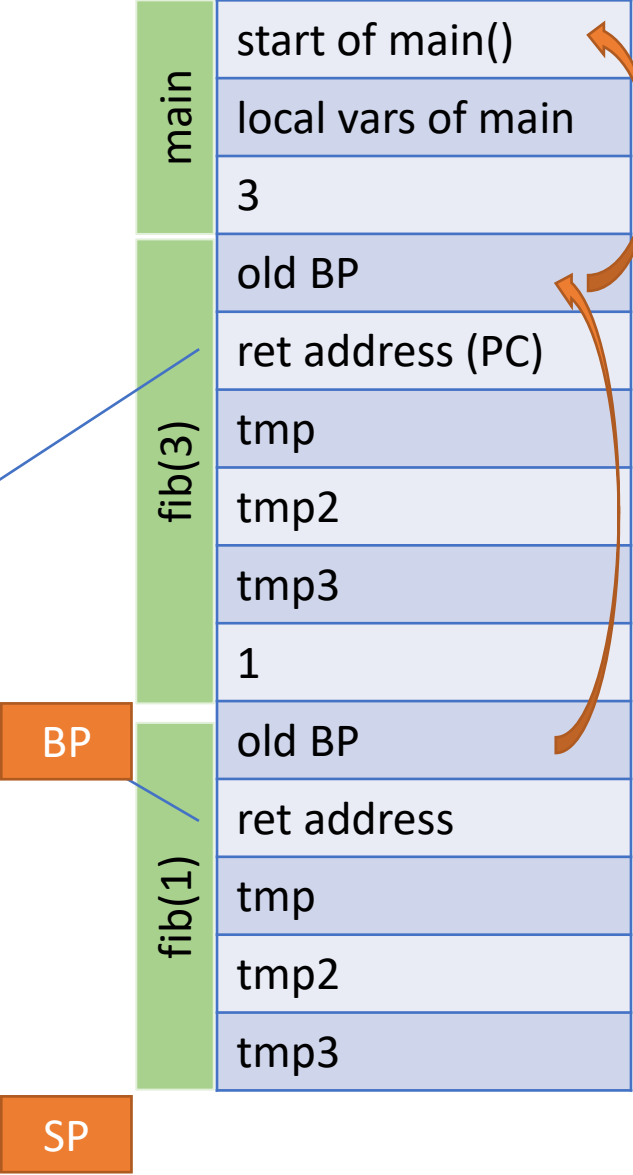
CX

DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    PC int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

fib(3)

Stack:



AX

1

BX

CX

DX

```
int fib(int n) {
    if (n == 0)
        return 0
    if (n == 1)
        return 1
    int tmp = fib(n - 2)
    int tmp2 = fib(n - 1)
    int tmp3 = tmp1 + tmp2
    return tmp3
}

fib(3)
```

PC

BP

SP

Stack:

main	start of main()
	local vars of main
	3
fib(3)	old BP
	ret address (PC)
	tmp
	tmp2
	tmp3
	1
	old BP
	ret address
	tmp
	tmp2
	tmp3

AX

1

BX

CX

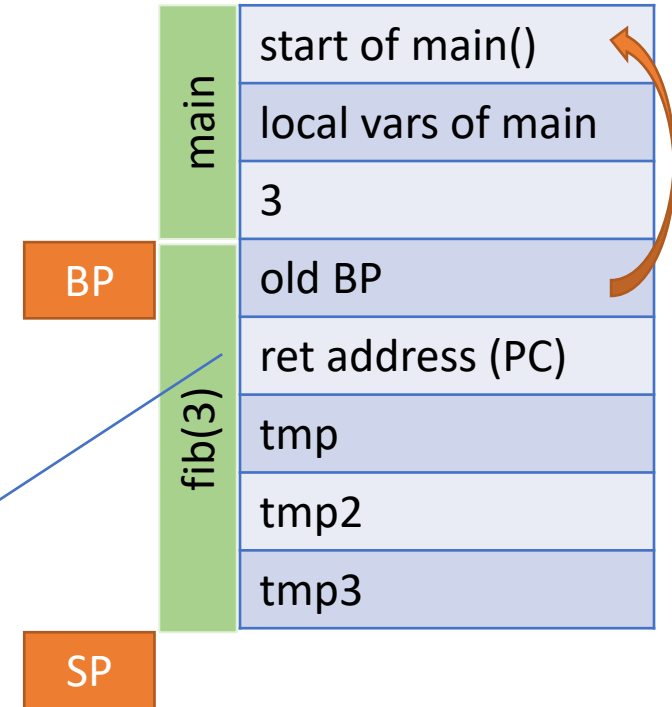
DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC

fib(3)

Stack:



AX
BX
CX
DX

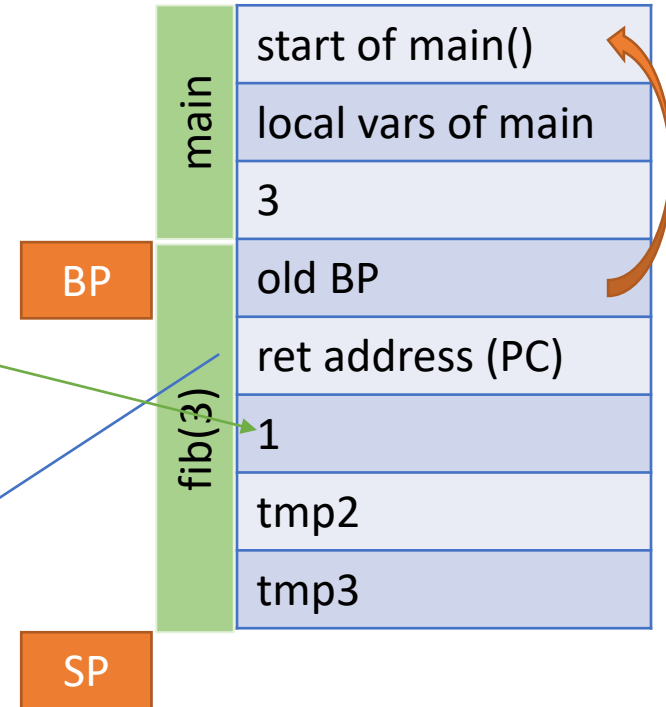
1

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC

fib(3)

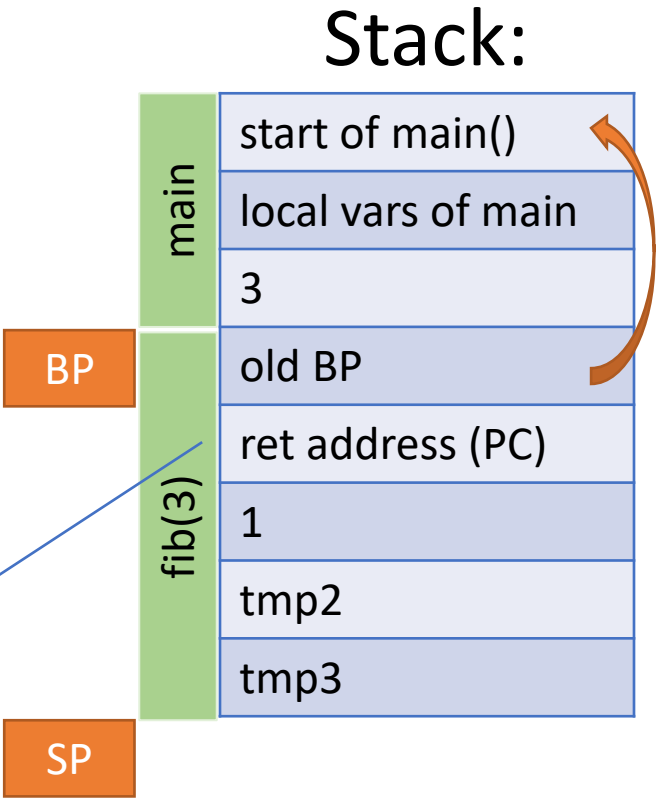
Stack:



AX
BX
CX
DX

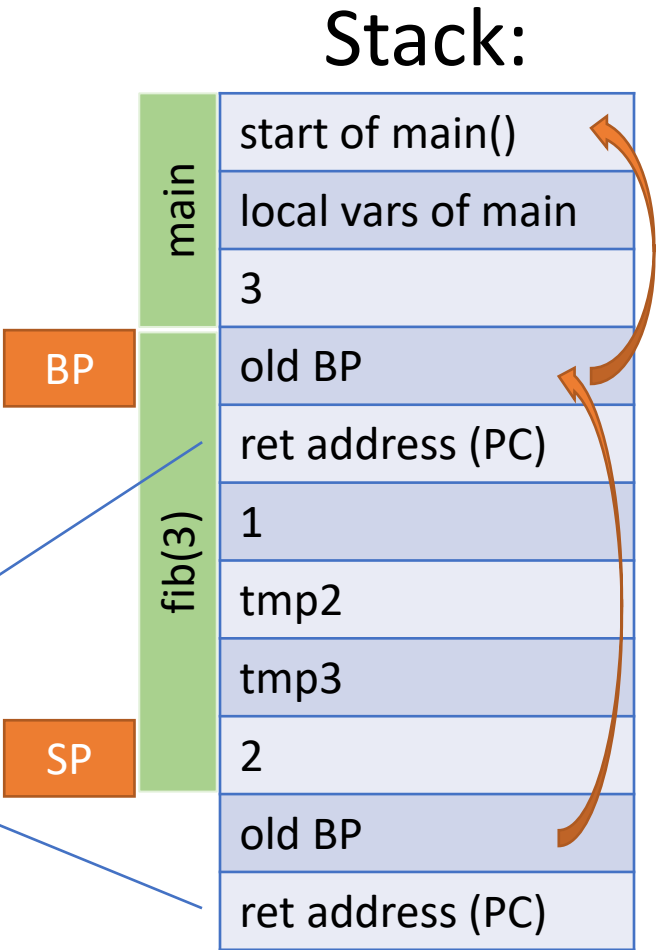
```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    PC int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

fib(3)



AX
BX
CX
DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    PC int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```

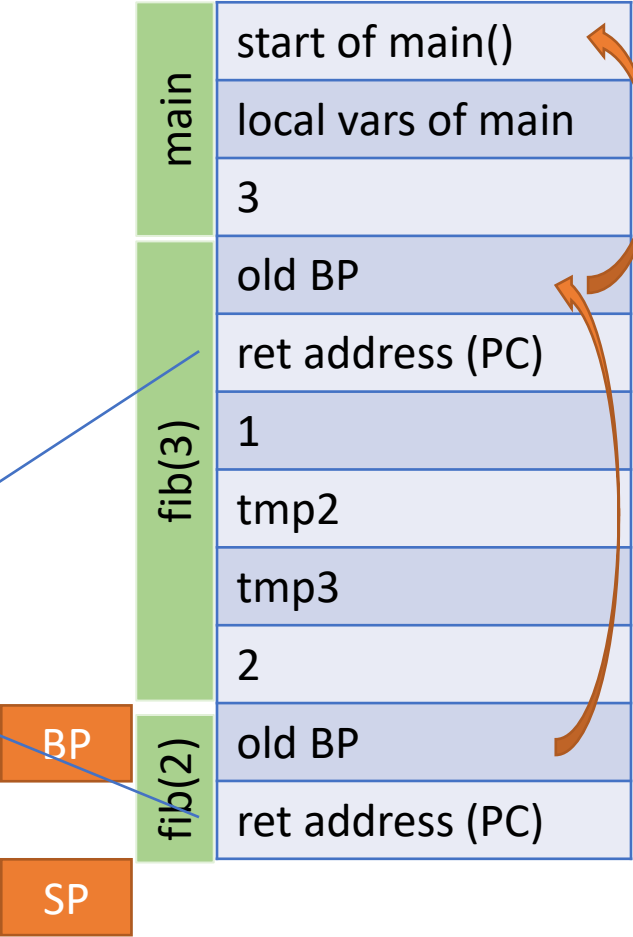


AX
BX
CX
DX

```
PC int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

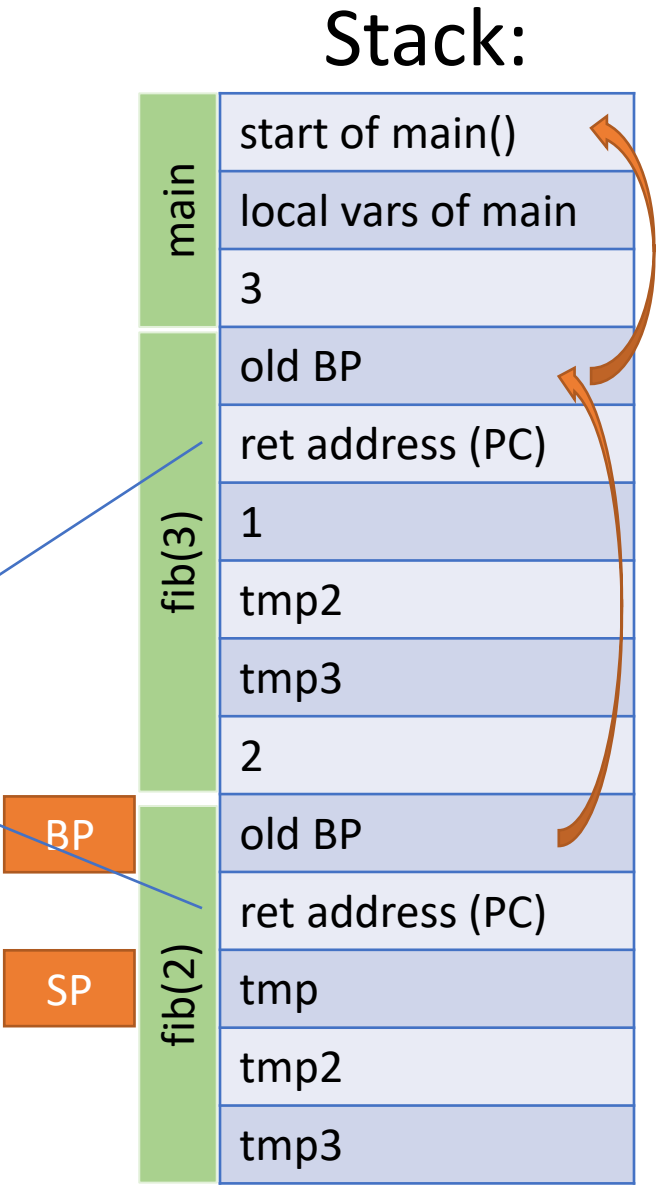
fib(3)

Stack:



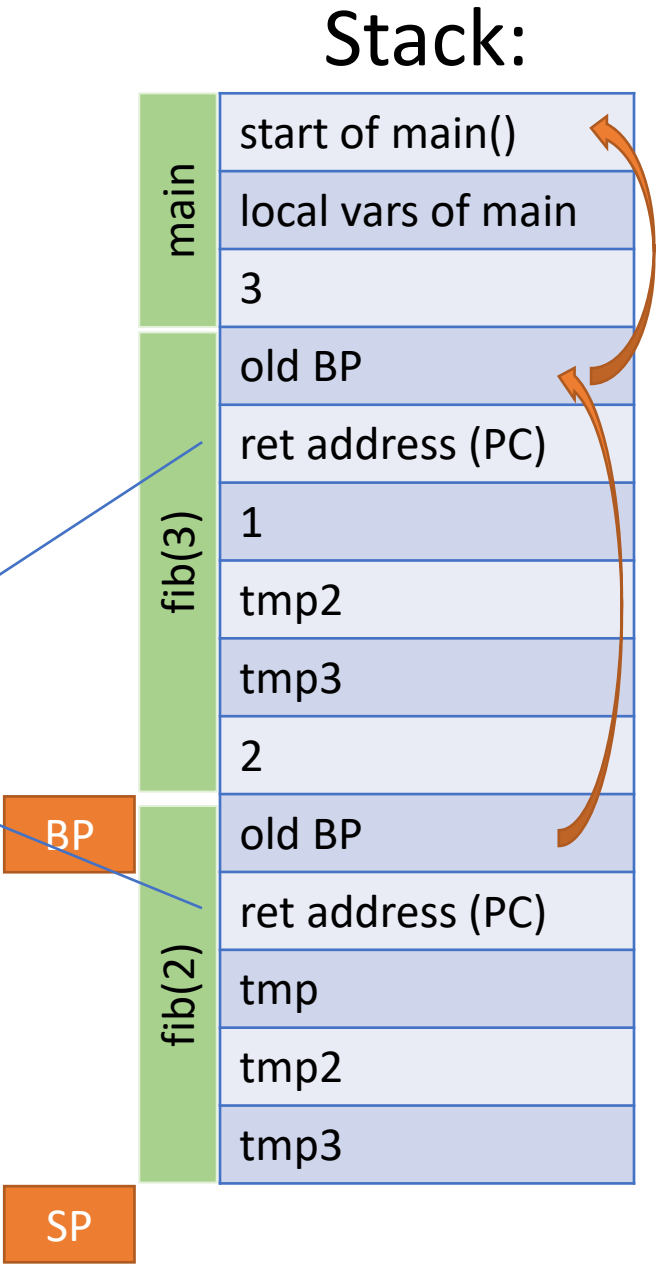
AX
BX
CX
DX

```
PC int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



AX
BX
CX
DX

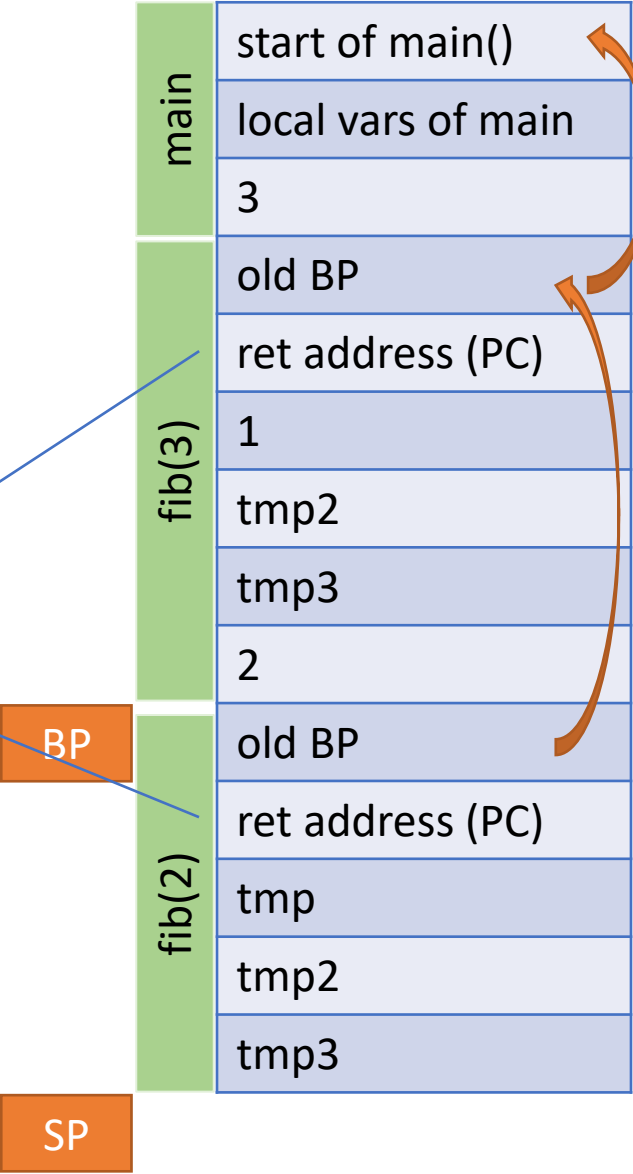
```
PC int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    PC int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```

Stack:



AX

BX

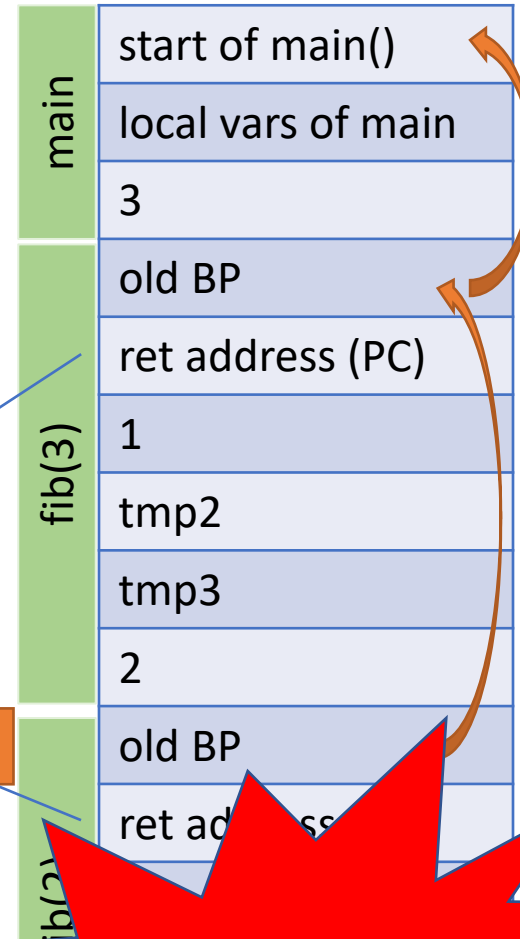
CX

DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    PC int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

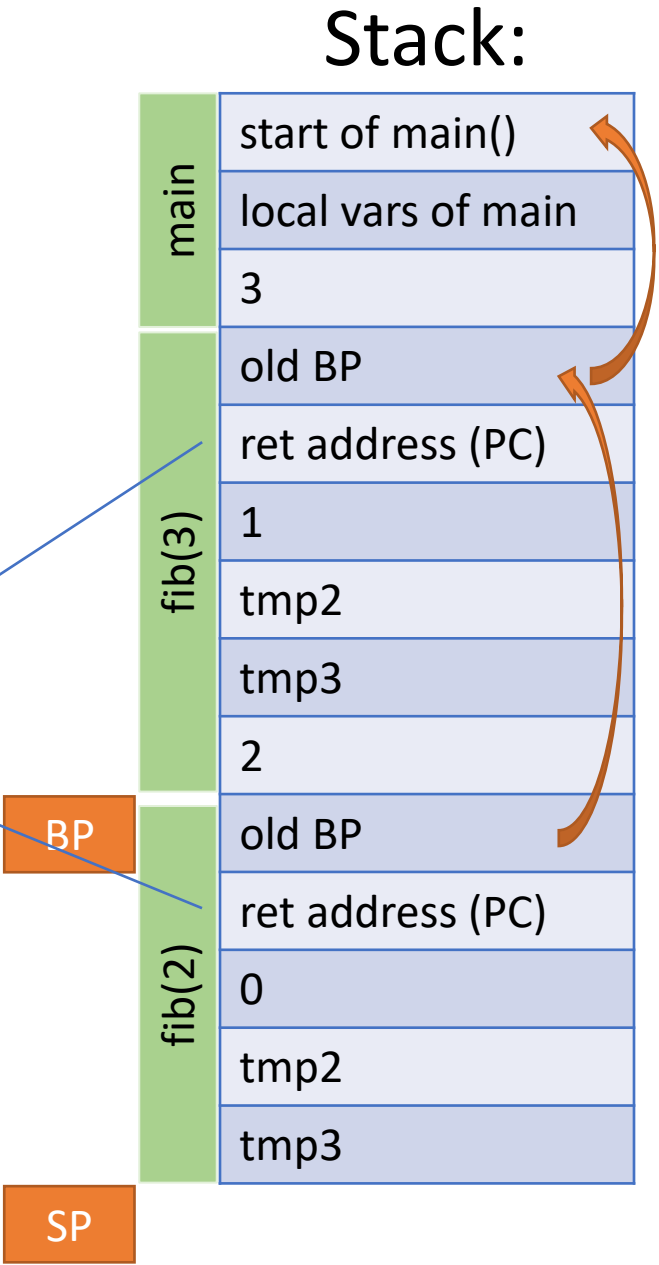
fib(3)

Stack:

Stack
Overflow

- AX
- BX
- CX
- DX

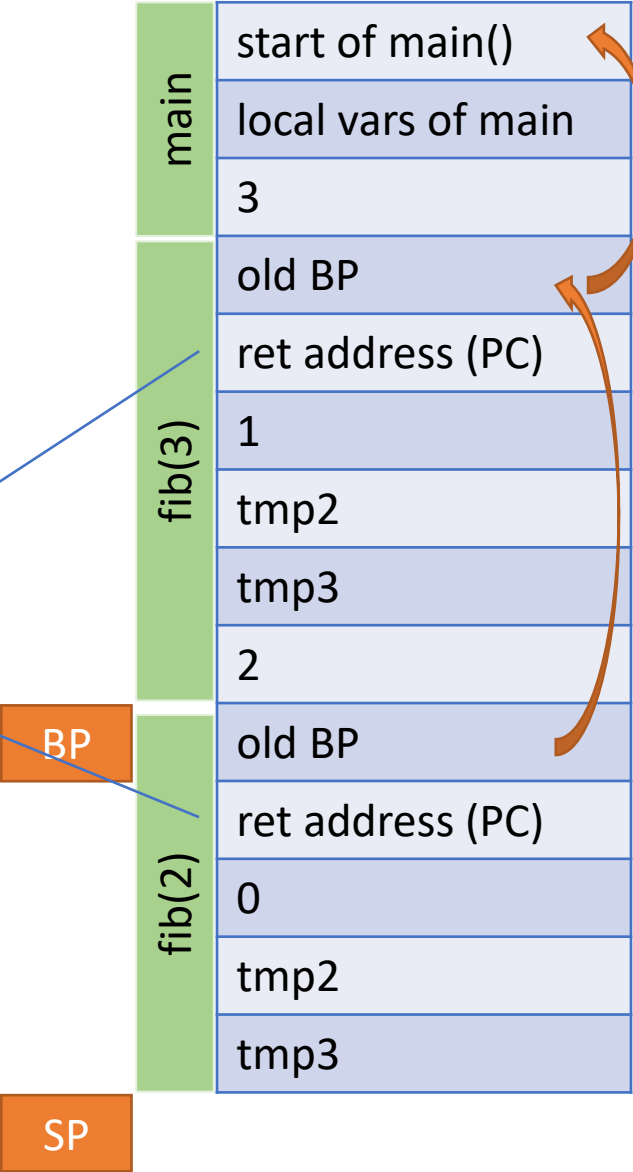
```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    PC int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    PC int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```

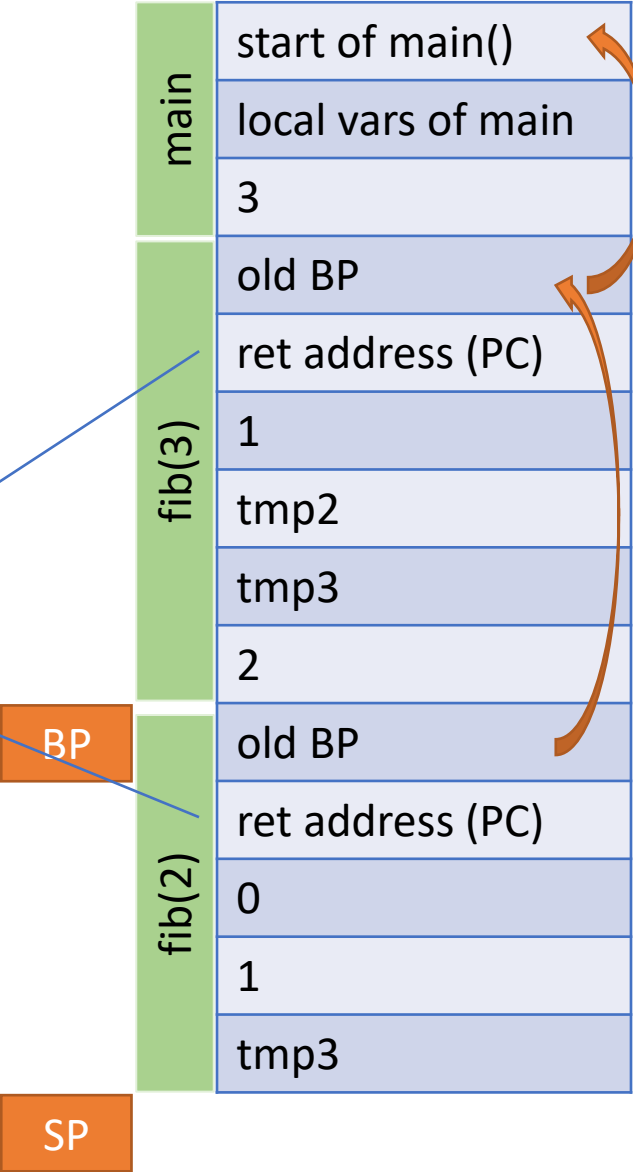
Stack:



- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    PC int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```

Stack:

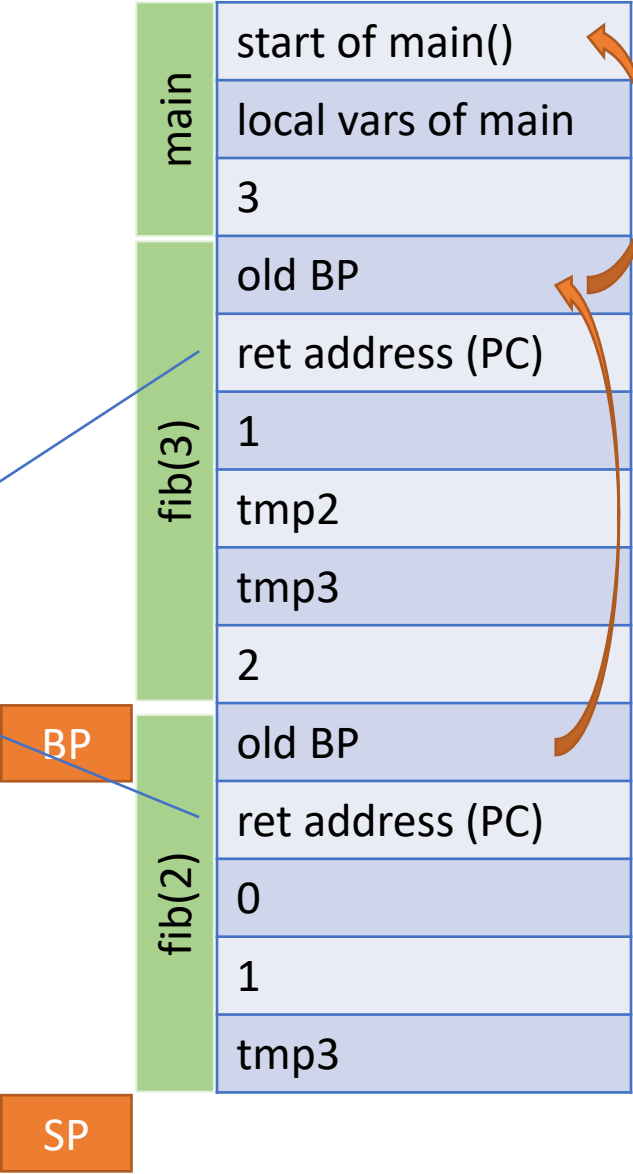


- AX
- BX
- CX
- DX

```
int fib(int n) {
    if (n == 0)
        return 0
    if (n == 1)
        return 1
    int tmp = fib(n - 2)
    int tmp2 = fib(n - 1)
    int tmp3 = tmp1 + tmp2
    return tmp3
}
```

fib(3)

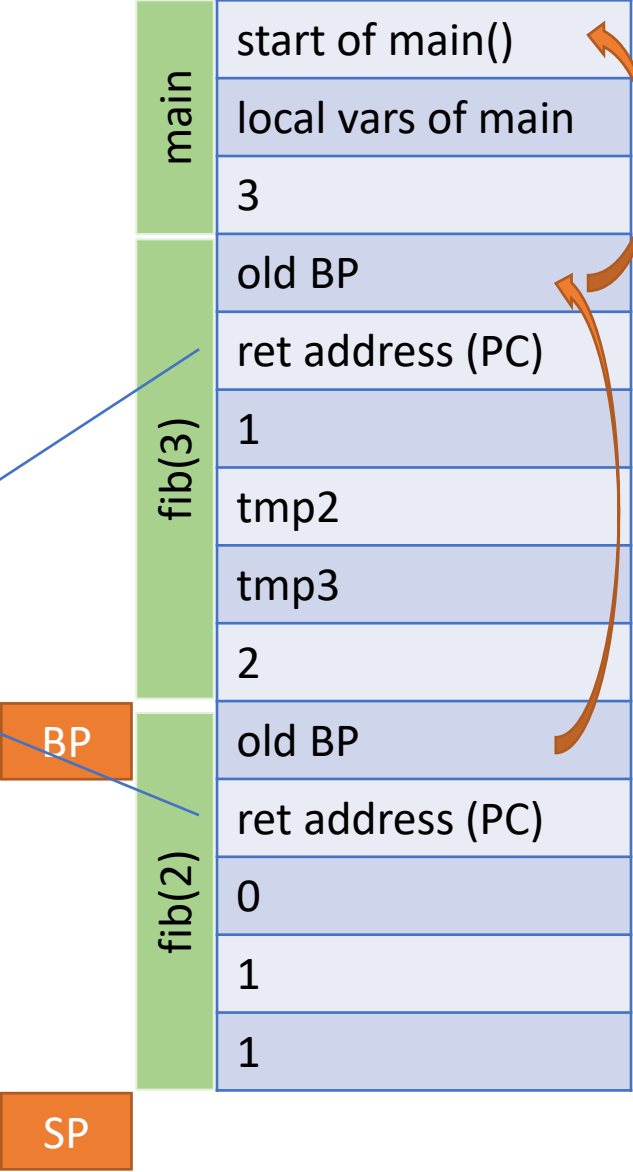
Stack:



- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    PC int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```

Stack:



AX
BX
CX
DX

1

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

fib(3)

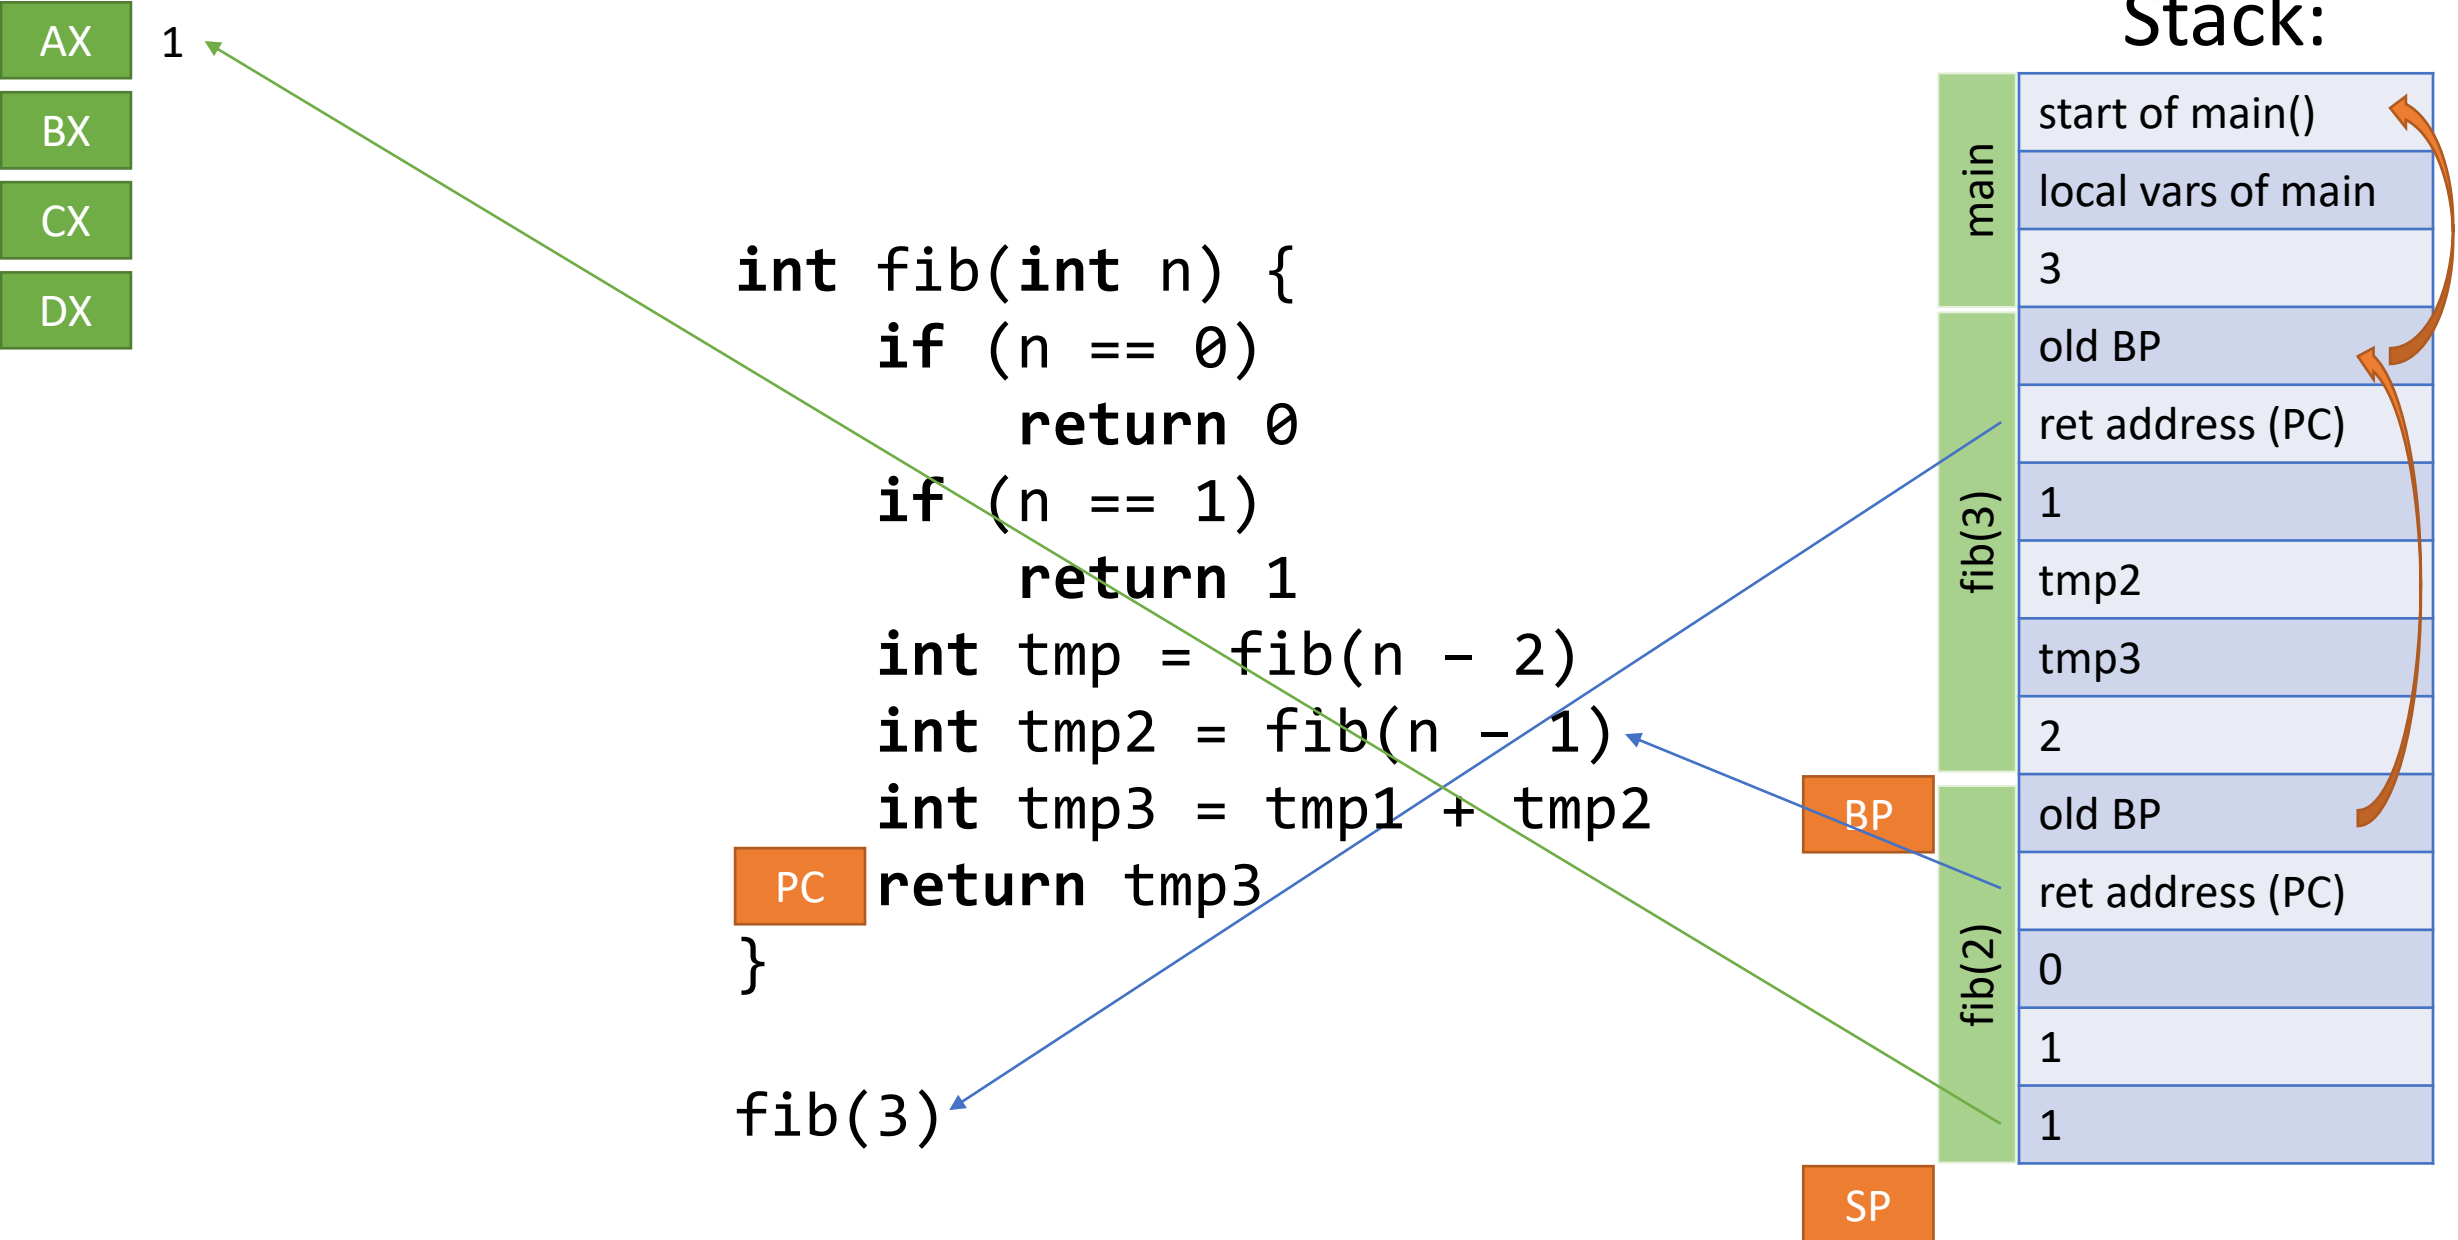
Stack:

main	start of main()
	local vars of main
	3
fib(3)	old BP
	ret address (PC)
	1
	tmp2
fib(2)	tmp3
	2
	old BP
	ret address (PC)
	0
	1
	1

BP

PC

SP



AX

1

BX

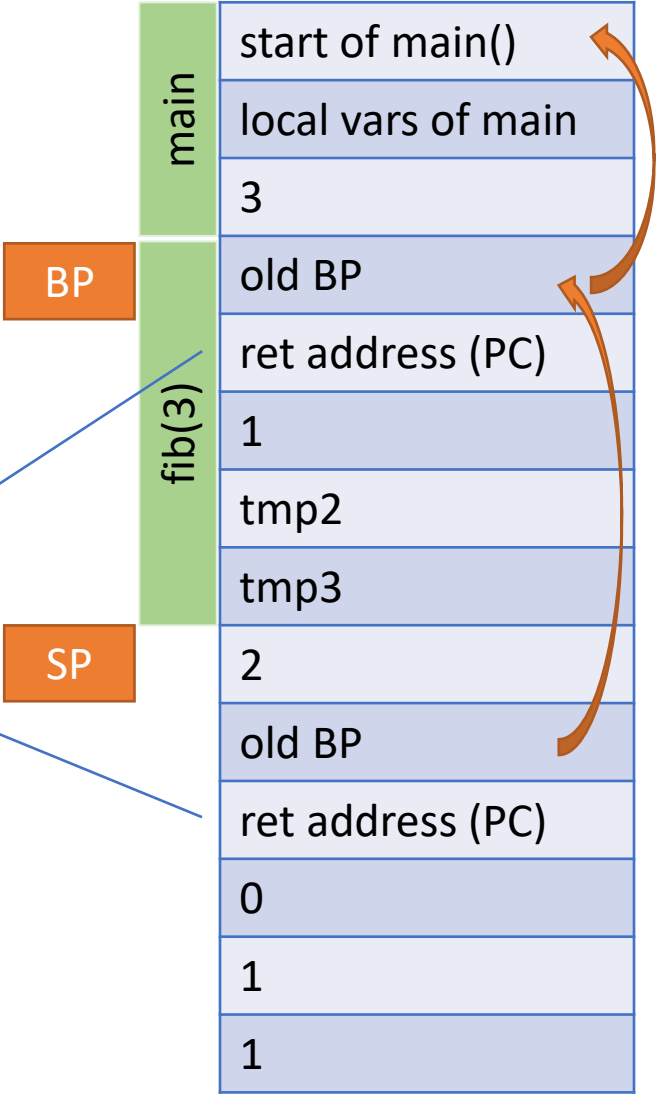
CX

DX

```
int fib(int n) {
    if (n == 0)
        return 0
    if (n == 1)
        return 1
    int tmp = fib(n - 2)
    PC int tmp2 = fib(n - 1)
    int tmp3 = tmp1 + tmp2
    return tmp3
}

fib(3)
```

Stack:






1

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC

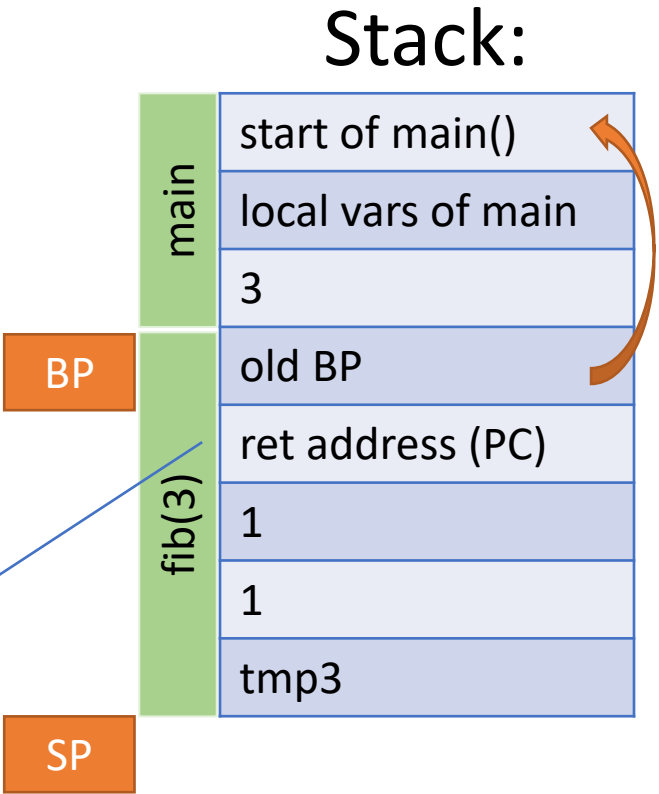
fib(3)

Stack:



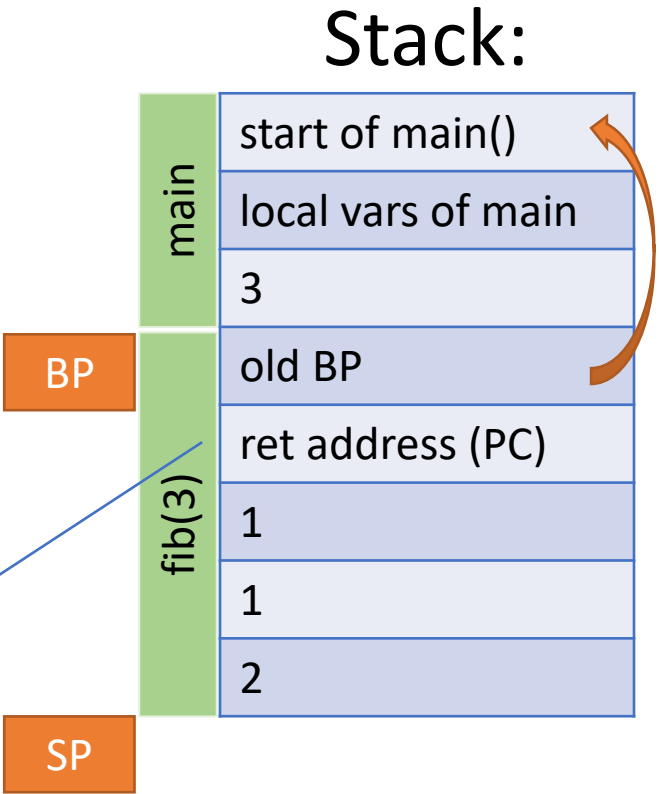
AX
BX
CX
DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    PC int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



AX
BX
CX
DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    PC int tmp3 = tmp1 + tmp2  
    return tmp3  
}  
  
fib(3)
```



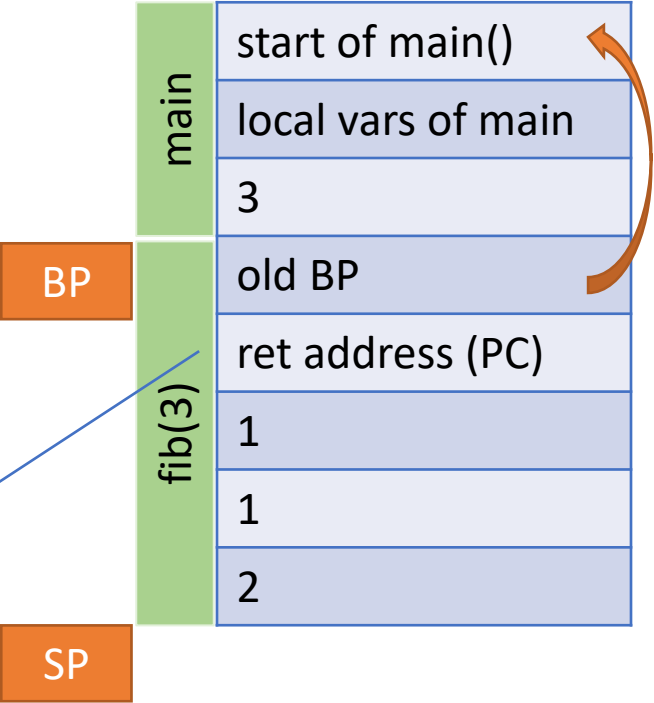
- AX
- BX
- CX
- DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC

fib(3)

Stack:



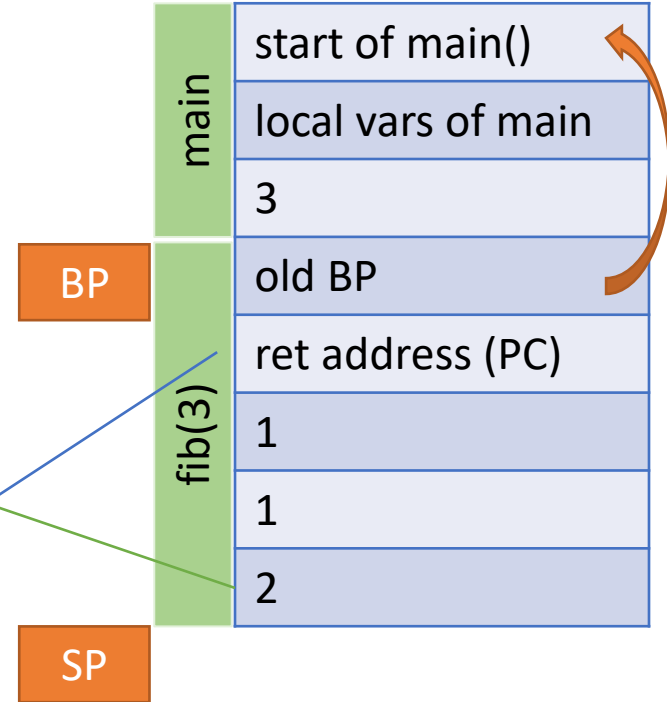


2

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

fib(3)

Stack:



AX

2

BX

CX

DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC fib(3)

Stack:

BP	main	start of main()
		local vars of main
SP		3
		old BP
		ret address (PC)
		1
		1
		2

AX

2

BX

CX

DX

```
int fib(int n) {  
    if (n == 0)  
        return 0  
    if (n == 1)  
        return 1  
    int tmp = fib(n - 2)  
    int tmp2 = fib(n - 1)  
    int tmp3 = tmp1 + tmp2  
    return tmp3  
}
```

PC

fib(3)

Stack:

BP


main

start of main()

local vars of main

SP

Call Stack

- stack frames used also for nested blocks
- elements on stack must have statically known sizes  why?
- slightly more complicated for languages that allow nested functions
 - stack frame saves the old BP as well as BP to the closest parent function
- when function returns data on stack cease to exist
 - this is **not** ok for other than local variables

Data Segment

- stores global variables
- has statically known size
- data in DS are valid throughout the duration of the program



Is this always true?

Data Segment

- stores global variables

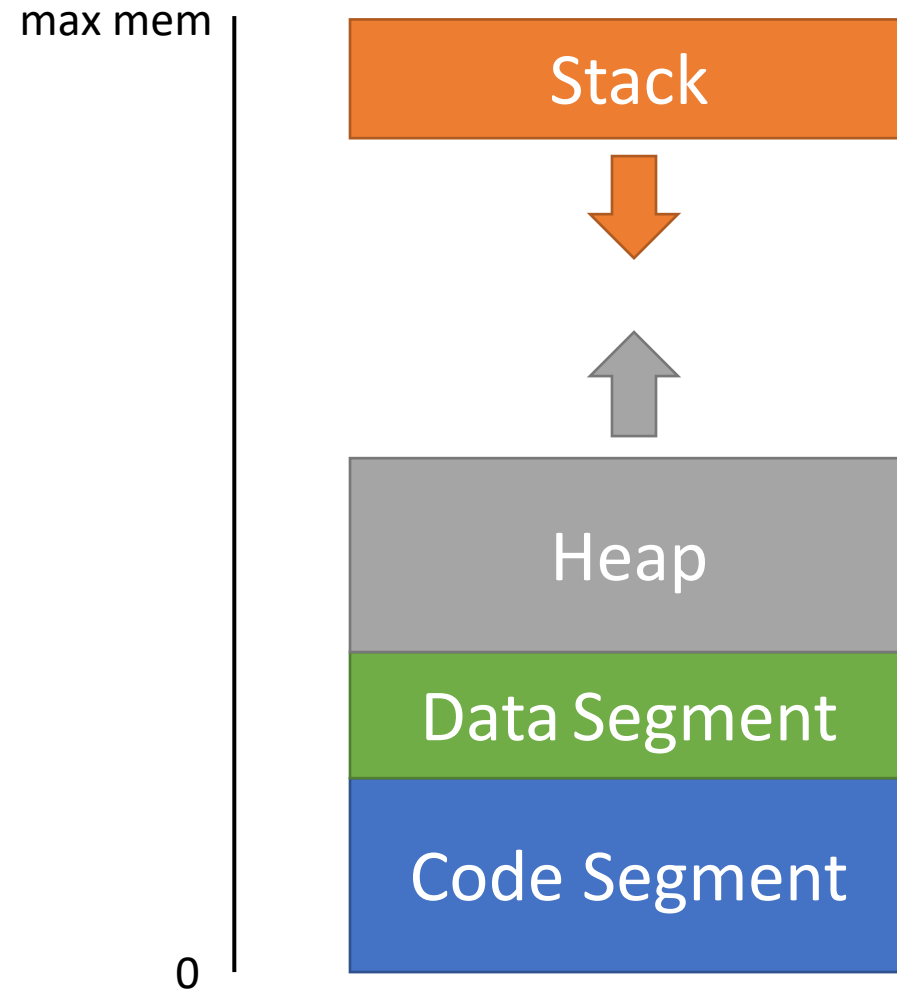
- has statically known size ← *But so does stack!!
Where to put dynamically sized data?*

- data in DS are valid throughout the duration of the program

Heap

- memory area for dynamically allocated data
- allows allocation & deallocation in arbitrary order
- manual or automatic (garbage collector) management
- dynamic size depending on the actual allocation at any given time (sort of)

Memory Layout



Calling Conventions


Calling Conventions

- how to pass arguments?
 - stack, registers, both, argument order
- how to return the result of the function?
 - stack, registers, type-dependent
- which registers are free for caller, and which for callee
- who is responsible for cleanup?
 - caller cleanup, callee cleanup
- architecture, OS, and language dependent

cdecl (x86)

- Microsoft's C compiler, caller saved
- arguments on stack, right to left
- result in eax, st0, or stack
- eax, ecx, edx caller saved, rest callee saved

cdecl (x86)



```
push    ebp        ; save caller's frame and start new frame
mov     ebp, esp
push    3           ; push the arguments left to right
push    2
push    1
call    f           ; call function f
add     esp, 12      ; remove arguments from frame (caller saved)
add     eax, 1       ; result returned in eax (if integer)
mov     esp, ebp    ; ebp is callee cleaned (call to itself!)
pop     ebp
ret
```

cdecl (x86)

- sometimes the bp is updated by callee (BP - addressing in callee, like our example)
- larger results returned on stack (caller allocated)
- results in more registers (eax:edx typically)
- stack alignment considerations

Other x86 Calling Conventions

- fastcall (arguments in registers, callee cleanup)
- thiscall (this pushed on stack as first argument, some variants in ecx, caller cleanup)
- stdcall (callee cleanup, right to left args, Win32API)
- x86_64 calling conventions (Microsoft vs the world)

ARM (A32)

- link register (LR) for return address
 - faster than stack for calls to leaf subroutines
 - not really that necessary with inlining
- more registers
 - r0, r1, r2, r3 = arguments (callee saved)
 - r4 – r11 = local variables (callee saved)

ARM64

- even more registers
- r19-r29 callee saved
- r9 – r15 caller saved
- r0-r7 arguments and results

Target Architectures

Target Architectures

- x86, x86-64, ARM, RISC-V, MIPS, Itanium, Sparc, AVR, ..., ...
- RISC, CISC, EPIC
- all have memory, registers, instructions

CISC

- memory is fast(!), and pricey
- clock speed is function of available technology only, cpu size & power is not an issue
- since there are transistors to spare, complex CPU instructions are possible and they greatly speed up the program and lower the memory requirements

CISC

- reg-mem arithmetics
- very complex addressing modes
- special instructions to support calls, control flow and other higher level language features
- variable instruction length

CISC

- memory is fast(!), and pricey
- clock speed is function of available technology only, cpu size & power is not an issue
- since there are transistors to spare, complex CPU instructions are possible and they greatly speed up the program and lower the memory requirements

CISC

- memory is *slow and cheap*
- clock speed is largely function of *cpu size & power*
- *cpu features on die compete with caches, pipeline stages, super-scalar ALUs, etc., there are no transistors to spare, complex instructions make the whole CPU slower*

RISC

- reg-reg arithmetics
- simple addressing modes
- fixed instruction length
- highly regular decoding

SuperScalar Processors

- as instruction is processed by the CPU, different parts are utilized
- instead of idling the unused circuits, more instructions can be processed at the same time
- instruction pipeline

Pipeline Stalls

- pipeline only helps if it is full
- branches, more complex ALU operations, data dependencies, memory accesses may cause holes (stalls) in the pipeline
- the instructions must be scheduled to minimize these

1980

RISC

- small, fast to execute instructions soon reached the limit of 1 IPC
- clock rates limited by technology & power
- can we go below 1 IPC?

VLIW

- Very Large Instruction Words
- large instructions encoding multiple smaller operations
- that are executed in parallel
- (the compiler schedules the operations into the VLIWs ahead of time)

EPIC

- non-determinism of memory access is a problem for VLIW
- Mitigated by EPIC (Explicitly Parallel Instruction Computing) where instructions and continuation logic is put by compiler into bundles
- explicit prefetching, explicit speculative pre-loading, predicated execution all used to make the bundles into more deterministic sequence

Intel Itanium

The Awesome Compiler

The Awesome Compiler

Not so much...

Super-Super Processors

- as superscalar or EPIC processors got more and more complex and powerful, the scheduling demands placed on the compiler were too great
- multiple execution units per stage were added to lower the congestion
- the super-super architecture

SMT

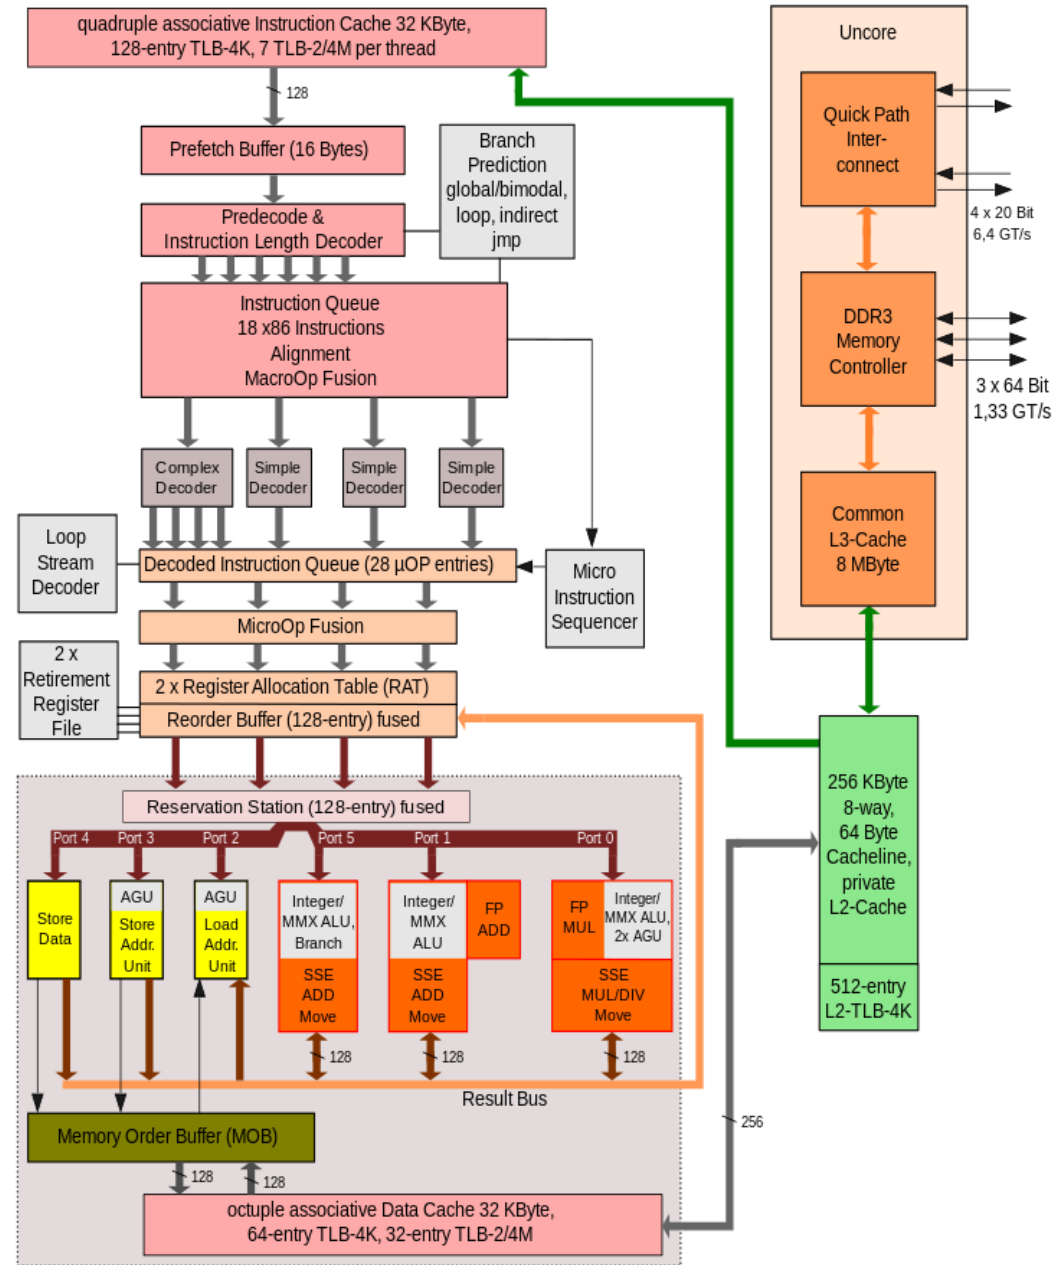
- often computers do more than one thing anyways
- SMT (Simultaneous Multithreading) makes multiple “cores” share parts of the execution units
- better utilization of the resources
- no thread context switching required

Hyperthreading

Out of Order Execution

- CPU keeps a pool of fetched instructions
- these are issued dynamically not according to their order in program, but availability of required execution units and inputs
- requires complex bookkeeping to preserve sequential semantics
- makes compiler scheduling much less important

Intel Nehalem microarchitecture



Practical ISA

```
struct pos {  
    int32_t x;  
    int32_t y;  
};  
pos items[128];
```

```
int32_t j = items[i].y;
```

```
// r1 contains & items  
// r2 contains i
```

```
ldi r3, r2
```

```
shr r3, 3
```

```
add r3, r1
```

```
ldi r4, 4
```

```
add r3, r4
```

```
load r4, r3
```

** by 8 (sizeof pos)*



load from memory



```
struct pos {  
    int32_t x;  
    int32_t y;  
};  
pos items[128];
```

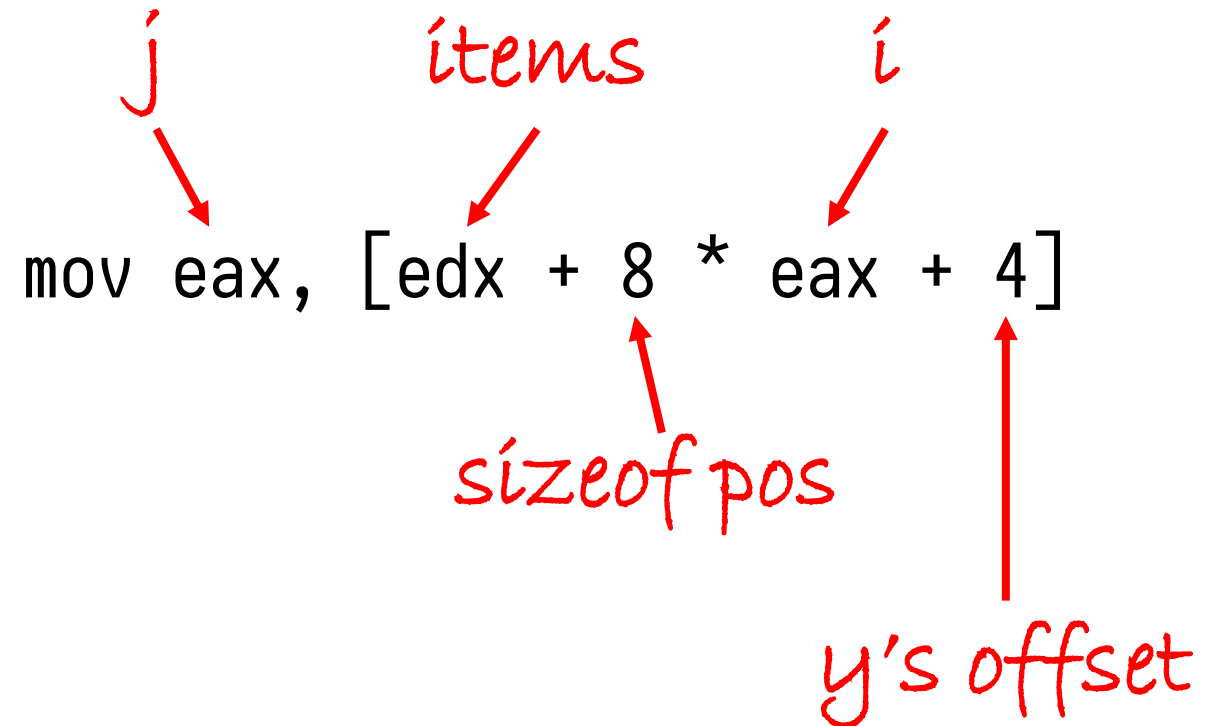
```
int32_t j = items[i].y;
```

```
// edx contains & items  
// eax contains i
```

j *items* *i*

```
mov eax, [edx + 8 * eax + 4]
```

sizeof pos *y's offset*




```
struct pos {  
    int32_t x;  
    int32_t y;  
};  
pos items[128];  
  
int32_t * j = & items[i].y;  
  
// r1 contains & items  
// r2 contains i  
  
ldi r3, r2  
shr r3, 3  
add r3, r1  
ldi r4, 4  
add r3, r4
```

*r3 contains the
address*



```
struct pos {  
    int32_t x;  
    int32_t y;  
};  
pos items[128];
```

```
int32_t * j = & items[i].y;  lea eax, [edx + 8 * eax + 4]
```

Fun with CISC

- complex instructions not created equal
- some are deprecated (made slower), or made faster
- often complex instructions can be used for surprising purposes
- consider lea

```
lea ecx, [edx + 8 * eax + 4]
```

```
lea ecx, [edx + 8 * eax + 4]
```

```
mov ecx, eax
```

```
shl ecx, 3
```

```
add ecx, 4
```

```
add ecx, edx
```

Code Size Matters

- variable length encodings even on RISC ISAs
- immediate argument sizes
- relative branch distances
- not all registers created equal

Further Reading

<https://medium.com/swlh/what-does-risc-and-cisc-mean-in-2020-7b4d42c9a9de>