

Code Generation

NI(E)-GEN, Spring 2021

<https://courses.fit.cvut.cz/NI-GEN>



FIT

```
int foobar(int a, int b) {  
    int c;  
    if (a < b)  
        c = a + b;  
    else  
        c = a - b;  
    return c;  
}
```

```

int foobar(int a, int b) {
    int c;
    if (a < b)
        c = a + b;
    else
        c = a - b;
    return c;
}

```

```

0 = ARG_ADDR 0 // &a
1 = ARG_ADDR 1 // &b
2 = ALLOC_L int // &c
3 = LOAD [0] // a
4 = LOAD [1] // b
5 = LT 3 4 // a < b ?
   COND_JUMP 5 ift iff

```

```

ift:
6 = LOAD [0] // a
7 = LOAD [1] // b
8 = ADD 6 7 // a + b
   STORE [2] 8 // c =
   JMP after

```

```

iff:
9 = LOAD [0] // a
10 = LOAD [1] // b
11 = SUB 9 10 // a - b
    STORE [2] 11 // c =
    JMP after

```

```

after:
12 = LOAD [2] // c
    RET 12

```

```

int foobar(int a, int b) {
    int c;
    if (a < b)
        c = a + b;
    else
        c = a - b;
    return c;
}

```

```

0 = ARG_ADDR 0 // &a
1 = ARG_ADDR 1 // &b
2 = ALLOC_L int // &c
3 = LOAD [0] // a
4 = LOAD [1] // b
5 = LT 3 4 // a < b ?
COND_JUMP 5 ift iff

```

ift:

```

6 = LOAD [0] // a
7 = LOAD [1] // b
8 = ADD 6 7 // a + b
STORE [2] 8 // c =
JMP after

```

iff:

```

9 = LOAD [0] // a
10 = LOAD [1] // b
11 = SUB 9 10 // a - b
STORE [2] 11 // c =
JMP after

```

after:

```

12 = LOAD [2] // c
RET 12

```

```

int foobar(int a, int b) {
    int c;
    if (a < b)
        c = a + b;
    else
        c = a - b;
    return c;
}

```

returns argument value

```

0 = ARG 0 // a
1 = ARG 1 // b
2 = LT 0 1 // a < b ?
    COND_JUMP 5 ift iff

```

ift:

```

3 = ADD 0 1 // a + b
    JMP after

```

iff:

```

4 = SUB 0 1 // a - b
    JMP after

```

after:

```

5 = PHI ift 3, iff 4
    RET 5

```

the usual phi node...

```
int foobar(int a, int b) {  
    int c;  
    if (a < b)  
        c = a + b;  
    else  
        c = a - b;  
    return c;  
}
```

```
0 = ARG 0 // a  
1 = ARG 1 // b  
2 = LT 0 1 // a < b ?  
    COND_JUMP 5 ift iff  
ift:  
3 = ADD 0 1 // a + b  
    RET 3  
iff:  
4 = SUB 9 10 // a - b  
    RET 4
```

Registers Are Not Perfect

- not all variables can be kept in registers
- some need addresses

```
int foobar(int a, int b) {  
    int c;  
    if (a < b)  
        c = a + b;  
    else  
        c = a - b;  
    return c;  
}
```



```
int foobar(int a, int b) {  
    int c;  
    if (a < b)  
        c = a + b;  
    else  
        c = a - b;  
    otherfoo(&c);  
    return c;  
}
```

Registers Are Not Perfect

- not all variables can be kept in registers
- some need addresses
 - references, pointers
 - shared across threads
 - ...
- some start with an address
 - heap allocated
 - special memory (ports, DMA, etc.)

IR

- unlimited registers
 - SSA (not necessary)
- all registers created equal
 - i.e. any register can be used in any situation

IR

- unlimited registers
 - SSA (not necessary)
- all registers created equal
 - i.e. any register can be used in any situation

Target

- only a few registers
- not all registers equal
 - certain instructions only work with certain registers

Register Allocation

Allocation & Assignment

- allocation determines which values will stay in registers and which won't
- assignment = for each value allocated to be in register, select the appropriate register

Value Spilling

- if there is not enough target registers, values must be stored in memory
- simple, remember we started with all values in memory in the IR
 - allocate space for the spilled values on stack
 - load / store as needed

Local Register Allocation

= *single basic block*

- in the beginning
 - all values in memory
- first read is load
 - all other reads just reuse the value in register
 - target registers, i.e. not SSA at this time
- last write is store
 - so that next basic blocks will have the information in memory

Local Register Allocation

- top-down:
 - loads & stores are costly, best savings when the most frequently accessed values are kept in registers
 - calculate usage frequencies, most frequent ones in registers
 - reserve enough registers for temporaries (statically known)

mov r0, 12

mov r1, 56

add r0, r1

add r1, 45

add r0, r1

add r0, 12

mov r3, r0

add r3, 12

```
mov r0, 12
mov r1, 56
add r0, r1
add r1, 45
add r0, r1
add r0, 12
mov r3, r0
add r3, 12
```

R0: 4 reads + 4 writes = 8

R1: 2 reads + 2 writes = 4

R3: 1 read + 2 writes = 3

```
mov r0, 12
mov r1, 56
add r0, r1
add r1, 45
add r0, r1
add r0, 12
mov r3, r0
add r3, 12
```

R0: 4 reads + 4 writes = 8

R1: 2 reads + 2 writes = 4

R3: 1 read + 2 writes = 3

```
mov r0, 12
mov r1, 56
add r0, r1
add r1, 45
add r0, r1
add r0, 12
mov r3, r0
add r3, 12
```

R0: 4 reads + 4 writes = 8

R1: 2 reads + 2 writes = 4

R3: 1 read + 2 writes = 3

we have no registers to write this since
all are occupied by values

Local Register Allocation

- top-down:
 - loads & stores are costly, best savings when the most frequently accessed values are kept in registers
 - calculate usage frequencies, most frequent ones in registers
 - reserve enough registers for temporaries (statically known)
- allocates registers for the **entire** basic block
 - this is often not necessary as values may not be used throughout the block
 - before first/after last use their register can be used for other values at no cost

Local Register Allocation

- bottom-up
 - at each instruction re-think the allocation
 - ensure operands are in registers
 - ensure result will be in register
 - if value no longer needed, forget register association
 - if there are no free registers and one is needed, use heuristic to spill
- heuristic
 - spill those that are used the furthest in the future

mov r0, 12

mov r1, 56

add r0, r1

add r1, 45

add r0, r1

add r0, 12

mov r3, r0

add r3, 12

free : ax bx

mov r0, 12

mov r1, 56

add r0, r1

add r1, 45

add r0, r1

add r0, 12

mov r3, r0

add r3, 12

free : bx

ax = r0

mov ax, 12

mov r0, 12

mov r1, 56

add r0, r1

add r1, 45

add r0, r1

add r0, 12

mov r3, r0

add r3, 12

free :

ax = r0

bx = r1

mov ax, 12

mov bx, 56

```
mov r0, 12
mov r1, 56
add r0, r1
add r1, 45
add r0, r1
add r0, 12
mov r3, r0
add r3, 12
```

free :

ax = r0

bx = r1

```
mov ax, 12
mov bx, 56
add ax, bx
add bx, 41
add ax, bx
```

mov r0, 12

mov r1, 56

add r0, r1

add r1, 45

add r0, r1

add r0, 12

mov r3, r0

add r3, 12

free : bx

ax = r0

mov ax, 12

mov bx, 56

add ax, bx

add bx, 41

add ax, bx

```
mov r0, 12
mov r1, 56
add r0, r1
add r1, 45
add r0, r1
add r0, 12
mov r3, r0
add r3, 12
```

free :

ax = r0

bx = r3

```
mov ax, 12
mov bx, 56
add ax, bx
add bx, 41
add ax, bx
add ax, 12
mov bx, ax
```

```
mov r0, 12
mov r1, 56
add r0, r1
add r1, 45
add r0, r1
add r0, 12
mov r3, r0
add r3, 12
```

```
free : ax
bx = r3
```

```
mov ax, 12
mov bx, 56
add ax, bx
add bx, 41
add ax, bx
add ax, 12
mov bx, ax
```

```
mov r0, 12
mov r1, 56
add r0, r1
add r1, 45
add r0, r1
add r0, 12
mov r3, r0
add r3, 12
```

```
free : ax
bx = r3
```

```
mov ax, 12
mov bx, 56
add ax, bx
add bx, 41
add ax, bx
add ax, 12
mov bx, ax
add bx, 12
```

Local Register Allocation

- bottom-up
- the devil lies in the details
 - don't allocate two registers for same value
 - reuse return register if arguments no longer necessary
 - tag clean/dirty register values, prefer spilling clean values (eliminates store)

This is all nice...but:

This is all nice... *but:*

- code is rarely a single basic block
 - and basic blocks tend to be short (not much to optimize)
- can be faster!
 - local techniques do not optimize reuse across basic blocks
- is actually very slow:
 - passing values between basic blocks still involves memory (expensive)
 - **all** touched values are passed (more memory accesses)
- we need to consider whole functions *hard: (*

Global Register Allocation

- must account for control flow
- operates on IR register live ranges (NI-APR)

Live Ranges

- a range from first to last use (load/store) of a variable

```
int foo(int a) {  
    int b = 1;  
    int d;  
    if (a == 5) {  
        b = a;  
    }  
    d = someFunction(1,3);  
    return b + d;  
}
```

Live Ranges

- a range from first to last use (load/store) of a variable

```
int foo(int a) {  
    int b = 1;  
    int d;  
    if (a == 5) {  
        b = a;  
    }  
    d = someFunction(1,3);  
    return b + d;  
}
```

Live Ranges

- a range from first to last use (load/store) of a variable

```
int foo(int a) {  
    int b = 1;  
    int d;  
    if (a == 5) {  
        b = a;  
    }  
    d = someFunction(1,3);  
    return b + d;  
}
```

a	b	
a	b	
a	b	
a	b	
	b	
	b	d
	b	d

Live Ranges

- a range from first to last use (load/store) of a variable

```
int foo(int a) {  
    int b = 1;  
    int d;  
    if (a == 5) {  
        b = 1;  
    } else {  
        b = a + 1;  
    }  
    d = someFunction(1,3);  
    return b + d;  
}
```

a	b	
a	b	
a	b	
	b	
	b	
a	b	
	b	
	b	d
	b	d

Global Register Allocation

- must account for control flow
- operates on IR register live ranges (NI-APR)

Linear Scan

- live ranges represented as intervals
- traversed chronologically, assigned registers in greedy manner
 - similar to bottom-up local approach
- fast (calculation) ...not so fast, actual code

Linear Scan

- live ranges represented as intervals
- traversed chronologically, assigned registers in greedy manner
 - similar to bottom-up local approach
- fast (calculation) ...not so fast, actual code
- the usual limitations of greedy algorithms
 - also, does not take holes in live ranges into account

So... Can we do better?

So... Can we do better?

If it's hard, make a graph of it...

- live ranges represented as graph nodes
- edges represent

- live ranges represented as graph nodes
- edges represent overlapping ranges

Register Allocation via Graph Coloring

- live ranges represented as graph nodes
- edges represent overlapping ranges
- colors for the graph represent target registers

is NP-complete :(

Graph Coloring

- key idea:
 - for K colors, a node with at most $K-1$ neighbors can always be safely colored. Just pick the color none of its neighbors have (1)
- find such nodes, color them and remove from the graph (2)
 - this might make some other nodes fall into the $< K$ neighbors category
- if no such nodes are found, spill a range and remove it
 - this reduces the degree of its neighbors, if there is one with $< K$ do (1), otherwise do (2)

Graph Coloring

- how to determine the color?
 - when removing nodes from graph, push to stack
 - when all on stack, pop from stack and assign lowest color none of its neighbors in the full graph has
- which ranges to spill?
 - this too complicates with control flow

Graph Coloring

- how to determine the color?
 - when removing nodes from graph, push to stack
 - when all on stack, pop from stack and assign lowest color none of its neighbors in the full graph has
- which ranges to spill?
 - this too complicates with control flow
 - assign spill costs to ranges (place/size/...)

The Other Things

- spilling can split the ranges
- ranges can be merged, if connected by a move
 - removes the move
- not all registers are created equal
 - and multiregister values
- whole program allocation
 - savings across function boundaries